Advanced Research Project

Zurich University of Applied Sciences

School of Management and Law

News-Driven Electricity Trading: Fine-Tuned NLP for Signal Generation in German Electricity Markets

---

Authors	Nicolas Alioth	(15-808-439)
	René Sunier	(20-630-638)
Program	MSs Banking & Finance	
Lecturer	Dr. Marc Weibel	
Department	Financial Data Science & Econometrics	

Submitted on: 11.01.2026

Management Summary

xxx

Table of Content

Table Reference V

List of illustrations VI

1 Introduction 1

1.1 Background 1

1.2 Research Question 4

1.3 Objective and Significance 4

1.4 Structure of the Research Project 5

2 Theoretical Foundations 6

2.1 Fundamentals of the Electricity Market 6

2.1.1 General Introduction to the Electricity Markets 6

2.1.2 Participants of the Electricity Markets 7

2.1.3 Instruments in the markets (forwards/next-day auctions) 8

2.1.4 Supply and Demand influences on the electricity markets 10

2.1.5 Motivation and common approaches to power trading 13

2.2 Fundamentals of Natural Language Processing 16

2.2.1 Introduction to NLP 16

2.2.2 Conceptual Foundations of Text-to-Quantitative Transformation 17

2.2.3 Limitations and Challenges of NLP in Financial Applications 20

2.3 Fundamentals of Trade Signal Generation using NLP 21

2.3.1 From Textual Information to Market Forecasts 22

2.3.2 From Forecasts to Trading Signals 25

3 Methodology 28

3.1 Framework Structure 28

3.2 Data Sourcing 28

3.2.1 News Data Collection 28

3.2.2 Energy Market Data Collection 30

3.3 Data Preprocessing 31
3.3.1 News Corpus Refinement 31
3.3.2 Energy Market Data Processing 33
3.4 Feature Engineering 34
3.4.1 Baseline Features 34
3.4.2 Topic Classification via Zero-Shot Learning 35
3.4.3 Sentence Embedding Generation 36
3.4.4 Time-Decay Aggregation of News Features 37
3.4.5 Dimensionality Reduction of Aggregated Embeddings 39
3.4.6 Feature Integration & Scaling 40
3.5 Model Training and Optimisation 41
3.5.1 Training Framework Architecture 41
3.5.2 Feature Parameter Selection via Ridge Regression 42
3.5.3 XGBoost Hyperparameter Optimisation 43
3.6 Model Evaluation Framework 48
3.7 Backtesting Framework 48
3.7.1 Trading Strategy Specification 48
4 Results 52
4.1 Data Overview and Descriptive Statistics 52
4.2 Feature Engineering Results 54
4.3 Predictive Performance and Statistical Evaluation 59
4.4 Economic Performance and Backtesting Results 62
5 Discussion 64
6 Conclusion 64
7 Bibliography 65
8 Appendix 65

Table Reference

Table 1: News Search Parameters 29
Table 2: Ridge Results 56
Table 3: LightGBM Classification Report (Test Set) 60
Table 4: Statistical Model Comparison 64

List of illustrations

Figure 1: Top 10 News Sources 53
Figure 2: News Coverage by Hour of Day 53
Figure 3: News Shocks and Price/Spread 54
Figure 4: Top 10 News Topic Classifications 55
Figure 5: News Embedding Visualization (UMAP) 56
Figure 6: Variation in Event Importance Over Time 58
Figure 7: XGBoost Top 10 Features (Gain) 59
Figure 8: Confusion Matrices for Signal and Baseline Model 60
Figure 9: Portfolio Equity Curve - Test Period 62
Figure 10: Drawdown (% of Peak) 63

List of Formulas

Es konnten keine Einträge für ein Abbildungsverzeichnis gefunden werden.

1 Introduction

The present chapter provides the conceptual and structural foundation for the study. The initial section (1.1) provides a contextual foundation, while the subsequent section (1.2) outlines the central research question. The objectives of the research project and its academic significance are addressed in Section 1.3, followed by an outline of the research project's structure in Section 1.4.

1.1 Background

Energy constitutes a fundamental input to modern economic activity and social welfare. From an economic perspective, energy enables production, transportation, and consumption across virtually all sectors, while from a societal standpoint it underpins living standards and technological development. Historically, macroeconomic growth has been closely associated with increasing energy consumption, reflecting the central role of energy in industrialisation and economic expansion (Stern, 2011, p. 26). In this sense, energy can be viewed not merely as a commodity, but as a prerequisite for organised economic life.

Within the broader concept of energy, electricity occupies a distinct and increasingly critical position. Unlike primary energy sources such as coal, gas, or oil, electricity represents a secondary form of energy that must be generated, transmitted, and consumed almost instantaneously. Its non-storability at scale implies that production and consumption must be continuously balanced, rendering electricity systems particularly sensitive to short-term shocks and forecast errors (Mack, 2014, p. 8). This characteristic differentiates electricity markets fundamentally from most other commodity markets.

The European power system has undergone substantial structural change over recent decades. Driven by climate policy objectives, technological progress, and regulatory reform, electricity generation in Europe has shifted increasingly towards renewable energy sources, particularly wind and solar power. While this transition has contributed to emissions reduction and diversification of supply, it has also increased the dependence of electricity generation on weather conditions, thereby amplifying short-term variability in supply (Dannecker, 2015, pp. 2–3). At the same time, geopolitical developments and energy security concerns have reinforced the importance of well-functioning power

markets and effective price signals within Europe (Mack, 2014, p. 105; European Council, 2025; IEA, 2025, pp. 41–42).

Electricity is produced and traded through a combination of physical generation assets and organised market mechanisms. Market participants can transact electricity via centralised exchanges, such as day-ahead and intraday markets, or through bilateral over-the-counter (OTC) agreements that allow for customised contracts and longer-term hedging (IEA, 2025, p. 18; Edwards, 2023, pp. 59–60). These market structures facilitate price discovery while accommodating the diverse objectives of generators, utilities, traders, and industrial consumers. However, the continuous balancing requirement and the growing share of intermittent generation introduce additional complexity into price formation.

Electricity prices are primarily determined by supply and demand conditions, including generation capacity, fuel costs, weather-driven load patterns, and network constraints. These factors are discussed in detail in Section 2.1. Beyond these structural drivers, prices are also influenced by information shocks that alter market participants' expectations. News related to weather forecasts, plant outages, grid congestion, regulatory interventions, or geopolitical developments can lead to rapid price adjustments as new information is incorporated into market valuations (Mack, 2014, pp. 105–106). In electricity markets, where uncertainty is high and adjustment mechanisms are immediate, such information effects are particularly pronounced.

The increasing relevance of news for price formation raises the question of how textual information can be systematically incorporated into market analysis and trading decisions. Traditional quantitative models typically rely on numerical inputs and are ill-equipped to process unstructured textual data at scale. Natural Language Processing (NLP) offers a methodological framework to address this limitation by transforming text into structured, quantitative representations that can be analysed using statistical and machine-learning techniques. In financial applications, NLP has been employed to extract sentiment, identify themes, and quantify narratives from news and reports, thereby capturing information that is not directly observable in market data (Das, 2014, pp. 3–5; Jurafsky & Martin, 2025, pp. 4–5).

A growing body of academic literature has examined how textual information can be exploited for forecasting and trading in financial and commodity markets, providing the

foundation for the present study. Early contributions focus on news sentiment and low-dimensional textual indicators, demonstrating that sentiment extracted from financial media contains predictive information for returns and volatility in commodity markets such as gold and oil (Smales, 2014; Li et al., 2021). These studies typically rely on aggregated sentiment measures derived from frequency-based representations and show that negative news often exerts a stronger market impact than positive news.

Subsequent research extends this line of inquiry by integrating text-derived features directly into forecasting models, treating textual information as a quantitative input alongside historical prices and fundamentals. Studies such as Guan et al. (2022) and Chakraborty et al. (2024) demonstrate that embedding-based representations of news significantly improve the prediction of commodity prices, particularly during periods characterised by geopolitical shocks or supply disruptions. This literature highlights the ability of NLP models to capture expectation-driven dynamics that precede observable market movements.

More recent contributions leverage transformer-based language models and large language embeddings, enabling richer semantic representations of news content. Research by Bai et al. (2024) and Bi et al. (2025) applies these advances explicitly to electricity markets, showing that news-augmented models outperform traditional price-, load-, and weather-based benchmarks, especially during periods of heightened volatility. These studies underline that textual information captures regulatory, infrastructural, and geopolitical developments that are not directly observable in physical system data.

Finally, a separate but related strand of literature addresses the translation of NLP-based forecasts into trading strategies, emphasising that predictive accuracy alone does not guarantee economic profitability. Studies across equity and commodity markets demonstrate that disciplined signal construction, confidence thresholds, and execution constraints are essential for transforming textual forecasts into viable trading decisions (Kirtac & Germano, 2024; J. Chen et al., 2025). This distinction between forecasting and tradable signal generation is particularly relevant in electricity markets, where price dynamics are fast, volatile, and subject to institutional constraints.

Against this background, the present study positions itself at the intersection of electricity market analysis and modern text analytics. By investigating how news-based information can be transformed into quantitative signals using NLP techniques, the paper seeks to

contribute to a deeper understanding of information processing in power markets and to explore the potential value of textual data for electricity price forecasting and trading strategies.

1.2 Research Question

This research project examines the following research question:

Can energy-related news be transformed into statistically and economically significant trading signals for the German power market using a dual Natural Language Processing framework based on transformer models and semantic embeddings?

To address this question, the study employs a structured modelling approach that combines quantitative market information with textual data. The baseline feature set comprises not only electricity price series but also broader power market indicators, such as baseload prices and temporal variables capturing intraday, weekly, and seasonal patterns. These features are augmented with NLP-based indicators extracted from energy-related news headlines, which are transformed into quantitative representations using transformer-based language models and embeddings. The integration of these feature sets enables the evaluation of whether news-derived information provides incremental predictive and economic value beyond traditional market-based inputs. Furthermore, the economic significance of these signals is assessed through a comprehensive backtesting framework, which evaluates the performance of the news-augmented signal model against both the baseline model and a naive long-only strategy. The detailed data construction, model architecture, and evaluation framework are presented in Chapter 3.

1.3 Objective and Significance

Electricity prices are highly sensitive to information that affects expectations about short-term supply and demand, including news on renewable generation, outages, weather conditions, policy interventions, and input costs such as natural gas, coal, and CO_2 allowances. While news-based forecasting and sentiment-driven trading strategies are well established in equity and commodity markets, comparatively little academic research has examined how electricity market news can be transformed into statistically and

economically meaningful trading signals. This gap is notable given the strong role of expectations and information shocks in short-term power markets.

The objective of this research project is to assess whether systematically incorporating energy-related news into electricity market models provides incremental value relative to approaches based solely on historical prices and temporal patterns. In particular, the study evaluates whether news-informed models improve both predictive accuracy and economic performance when applied to short-term trading decisions in the German power market.

The study contributes to the literature by extending recent advances in NLP and embedding-based text analysis to electricity markets and by explicitly distinguishing between forecasting performance and tradable signal construction. From a practical perspective, the results provide insight into whether structured processing of news can enhance short-term decision-making in power markets, where rapid information assimilation and disciplined execution are essential.

1.4 Structure of the Research Project

This research project is structured as follows. Chapter 2 presents the theoretical foundations and related literature, covering the fundamentals of electricity markets, key drivers of price formation, and the conceptual foundations of Natural Language Processing and text-based trading signal generation. This chapter establishes the economic and methodological context for the empirical analysis. Chapter 3 outlines the data sources, modelling framework, and methodological pipeline employed in the study. It describes the construction of market-based and news-derived features, the NLP models used to process textual information, and the procedures for model training, validation, and evaluation. Chapter 4 presents the empirical results of the forecasting and trading analyses. Model performance is evaluated using both statistical metrics and economic criteria, and the result of the news-augmented signal model is compared against the baseline benchmarks without news information and a naive long-only strategy. Chapter 5 concludes the research project by summarising the main findings, discussing limitations, and outlining implications for future research and practical applications in electricity markets. Chapter 6 provides a conclusion to the entire research project.

2 Theoretical Foundations

This chapter reviews academic literature and theoretical foundations of electricity price formation, news-based forecasting and the application of natural language processing in financial and energy markets. The analysis begins in Section 2.1, which provides a general overview of the power market, its institutional structure, and the fundamental drivers of electricity prices. Section 2.2 then outlines the fundamental concepts and methodological foundations of NLP, while Section 2.3 surveys extant research on news-based forecasting and trading applications in financial and commodity markets. Together, these sections establish the conceptual foundation for the empirical methodology developed in subsequent chapters.

2.1 Fundamentals of the Electricity Market

This section examines the structural and economic characteristics of electricity markets relevant to short-term price formation. The text provides an overview of the unique physical properties of power (Section 2.1.1), the roles of key market participants (Section 2.1.2), and the mechanisms of power trading (Section 2.1.3). Moreover, the section undertakes an analysis of the fundamental drivers of supply and demand (Section 2.1.4) and the subsequent trading behaviours (Section 2.1.5)

2.1.1 General Introduction to the Electricity Markets

The evolution of electricity systems has been characterised by the development of energy sources, with periods of both stagnation and rapid advancement. In recent times, there has been a resurgence in the adoption of renewable energy sources, marking a modern renaissance in this field. The 1970s oil crisis led to accelerated policy and technological shifts, which resulted in the revival of non-fossil energy sources. Today's power systems increasingly rely on a diverse mix of energy sources, including renewables like wind and solar power, along with established hydropower (Lefort & Simic, 2024). The underlying principle of renewable sources is that they are unlimited and not finite, in contrast to fossil fuels, which are utilised for conventional energy generation (Bundesnetzagentur, 2025a). As demonstrated by Statista (2025), the German market has experienced a notable increase in renewable electricity production, with wind power witnessing a 36% growth

and solar power recording a substantial 64% increase from 2018 to 2025. This expansion plays an important role in reducing emissions from electricity generation. However, as Dannecker (2015, p. 17) illustrates in the German context, it also brings certain challenges and unintended side effects. Wind and solar energy production is highly dependent on weather conditions, resulting in inconsistent output and impacting electricity prices, which leads to increased volatility. Countries that want to rely heavily on renewables, such as Germany, are particularly affected by these fluctuating energy prices. A more recent example of this phenomenon is the recent blackout that occurred in Spain this year, which was attributed to an excess of renewable energy generation and a deficit of conventional energy (CNN, 2025). In this context, Dannecker (2015, p. 20) highlights the need for more advanced forecasting methods for renewable electricity, ensuring that production can be better aligned with demand and helping to stabilise the increased volatility observed in modern power markets.

2.1.2 Participants of the Electricity Markets

Electricity generated at power plants is fed into the high-voltage transmission grid, where it can be transported over long distances with minimal losses. Transmission system operators (TSOs) manage this high-voltage network by ensuring that electricity from various generation facilities is safely integrated, routed, and delivered to regional distribution grids. At the European level, TSOs are coordinated by European Network of Transmission System Operators for Electricity (ENTSO-E), which harmonises grid operation, security standards, and cross-border electricity flows to support a reliable and interconnected power system (Bundesnetzagentur, 2025c). The transmission grid is the fundamental component of the electricity supply system, while the distribution grids pass electricity onward to households, firms, and smaller industrial consumers. Market participants purchase electricity on exchanges such as the European Power Exchange (EPEX) to either supply customers, hedge their exposures, or optimise their own generation portfolios (Bundesnetzagentur, 2025b). A notable issue is the relative immaturity of electricity as a commodity, in contrast to well-established products such as gold, coffee, or oil, which can be stored easily. Electricity, by contrast, must be consumed immediately, creating distinct challenges for its storage, pricing, and distribution (Edwards, 2023, p. 2).

2.1.3 Instruments in the markets (forwards/next-day auctions)

Power producers agree to sell electricity to buyers at pre-agreed terms via purchase agreements or long-term contracts, often structured physically or financially. These instruments are used to secure price certainty, facilitating financing, or, in some cases, to obtain exposure to market price fluctuations for speculative purposes (Edwards, 2023, p. 2). This can be achieved through bilateral OTC contracts or by trading through an intermediary such as EPEX, where transactions can be executed in the spot market or via derivatives (Edwards, 2023, p. 5). The present study focuses on spot and forward trading, while other contract forms are included for the sake of completeness, but fall outside the study's scope.

The two electricity markets under discussion in this research paper comprise the day-ahead auction and the intraday market, where they enable short-term balancing of supply and demand. The intraday market allows continuous trading, enabling participants to adjust positions in response to updated forecasts or unexpected changes in generation and consumption (EPEX, 2025). A distinctive feature of the spot market is that purchased electricity must be consumed immediately, which requires direct physical delivery to the end user. As a result, the spot market is accessible primarily to companies that are able to either take delivery or utilise the electricity directly. It is important to recognise that the spot market is characterised by pronounced volatility, driven by a combination of internal and exogenous factors. Unexpected events, such as unplanned outages, can lead to sharp price increases as the market reacts to sudden shifts in supply (Edwards, 2023, p. 51). Conversely, electricity prices may also become negative – an uncommon occurrence in commodity markets – when excess generation exceeds system demand and available flexibility. Negative prices therefore function as a market signal to curb inflexible generation and, where possible, stimulate additional consumption, helping the system absorb temporary oversupply arising from high renewable feed-in and low demand (EPEX, 2025).

Building on the characteristics of the spot market – where prices can fluctuate sharply due to volatility, scarcity, or even negative pricing – it is also important to consider how market participants manage these dynamics through longer-term instruments. In the day-ahead auction, bids and offers for each hourly delivery period of the following day are

aggregated into a uniform market-clearing price once the auction closes at 12:00 CET (EPEX, 2025). Beyond these hourly products, electricity is also traded using standardised delivery profiles. A base load contract represents all 24 hours of a given delivery period – such as a day, week, month, or year – and its price corresponds to the average of all hourly prices within that interval. A peak load contract, by contrast, covers only the hours with the highest typical demand, usually from 08:00 to 20:00 on working days, with its price reflecting the average of the hourly prices during this peak window, which are generally higher due to increased consumption (Müsgens & Bade, 2024, pp. 15–16).

Beyond these standardised delivery profiles, electricity markets also employ longer-term contractual instruments that allow price risks to be managed across extended horizons. Forward contracts are typically negotiated bilaterally over the counter and enable buyers and sellers to fix prices for future delivery periods, offering a high degree of flexibility with respect to maturity, volume and settlement structure (IEA, 2025, p. 62). Closely related to forwards, electricity swaps represent financially settled agreements that exchange fixed and floating price components and are widely used for risk management purposes. Complementing these bilateral instruments, futures contracts provide a standardised alternative that is traded on exchanges and cleared centrally, implying daily margining and reduced counterparty risk (IEA, 2025, p. 62; Mack, 2014, p. 52). Options further extend the set of available instruments by granting the right, but not the obligation, to buy or sell electricity or related underlyings, thereby allowing market participants to limit downside risk while retaining exposure to favourable price movements.

The coexistence of different delivery profiles, such as base and peak load, together with longer-term contractual instruments highlights the need to accommodate fluctuating electricity consumption and evolving risk exposures across time. While delivery profile's structure when electricity is supplied, forward and derivative products enable price risks associated with these profiles to be managed over extended horizons. Taken together, these instruments provide the foundation for organising trading activity in electricity markets and for translating physical system characteristics into market prices. Against this background, it becomes essential to consider the internal and external factors that influence price formation and shape the functioning of electricity markets more broadly.

2.1.4 Supply and Demand influences on the electricity markets

The existence of different delivery profiles, such as base and peak load, alongside the products traded in the spot and day-ahead markets, reflects the fact that electricity consumption does not remain constant throughout the day. These products allow market participants to address varying usage patterns and to handle the price movements that result from them. Against this background, it becomes essential to consider the supply and demand factors that shape price formation and determine how the electricity market functions more broadly. Although these factors are often conceptually separated, they are closely interlinked in practice, which means that a strict distinction is not always possible.

Previous sections have already highlighted demand-related factors such as base and peak load, which reflect the fact that electricity usage varies across the hours of the day. In addition, weekly cycles can also be observed, with noticeable differences between weekdays and weekends. These variations are largely driven by the recurring consumption behaviour of households and firms. Generators attempt to align production with these demand profiles as closely as possible, relying on forecasting models that are calibrated to expected usage patterns (Edwards, 2023, p. 135). Beyond daily and weekly patterns, electricity markets are also shaped by broader seasonal effects. The contrast between summer and winter plays a more pronounced role than in markets for storable commodities. For many of these commodities, storage creates a convenience yield – an economic benefit of holding inventories – that often contributes to a smoother, upward-sloping term structure. Many energy markets, including gas and electricity, by contrast, cannot rely on economically viable large-scale storage, and their prices therefore tend to exhibit more distinct seasonal patterns, with higher levels in winter and lower levels in summer due to the corresponding shifts in demand (Edwards, 2023, p. 77). Edwards (2023, p. 137) further shows a strong correlation between temperature and load. While electricity demand increases in winter due to heating requirements, hot summer days can likewise result in elevated consumption as cooling systems operate more intensively, thereby increasing the need for additional electricity supply.

Beyond these structural demand patterns, electricity markets are also influenced by broader macroeconomic conditions. Economic growth, industrial activity and business cycles affect electricity consumption through changes in production levels, investment activity and household demand, linking electricity prices to the general macroeconomic environment. Periods of economic expansion tend to raise electricity demand and tighten

capacity margins, while economic slowdowns reduce load growth and ease price pressure (Edwards, 2023, pp. 51 & 114).

As with demand, the supply side of the electricity market is shaped by a combination of technological, operational, meteorological and regulatory factors, which together determine both the maximum available generation and the marginal cost at which it can be supplied.

A central supply-side driver is the cost of fuels and other inputs, since these determine the marginal cost of conventional generation and thus the order in which power plants are dispatched. In economic terms, the relevant benchmark is often the levelized cost of electricity, which aggregates initial investment, fuel, operation and maintenance, and financing costs into a single unit cost over the lifetime of a plant (Mack, 2014, p. 212). For fossil-fuel technologies such as coal and gas, a large share of this cost is driven by fuel prices, so changes in coal, gas or CO_2 allowance prices feed directly into the short-run marginal cost of generation (Edwards, 2023, p. 138). By contrast, wind and solar are characterised by high upfront investment to build the plant but very low marginal costs, as additional output does not require extra fuel buying (IEA, 2025, pp. 24–25). This difference in cost structure explains why, once built, renewable plants typically enter the merit order before conventional units and push higher-cost generators to the margin. The merit order is a ranking of available power plants from lowest to highest marginal cost, ensuring that the cheapest energy sources are always used first to meet electricity demand. (Müsgens & Bade, 2024, pp. 6 & 96; Edwards, 2023, p. 131).

Beyond cost considerations, the technical characteristics of generating units play an important role in determining available supply. Conventional plants differ in their ramp-up and ramp-down capabilities, minimum stable generation levels and start-up times, which constrain their ability to respond quickly to variations in demand or renewable production (Edwards, 2023, p. 28). Intermittent renewable sources such as wind and solar exhibit strong weather-related fluctuations and relatively low average availability compared with their installed capacity. For example, analyses of wind power in Germany demonstrates that only a limited portion of the installed fleet can be considered reliably available at all times, with load factors spanning a wide range between very low and moderate utilisation levels (Dannecker, 2015, pp. 18–19). Edwards (2023, p. 358) likewise notes that wind output is inherently variable and sensitive to seasonal,

geographical and day time conditions. In addition to these technical and meteorological constraints, not all installed capacity is continuously usable. Parts of the generation fleet are periodically unavailable due to scheduled maintenance or unexpected outages. For this reason, adequacy assessments distinguish between unavailable and reliably available capacity, with system operators also maintaining reserve capacity to manage forecast errors and short-term imbalances (Dannecker, 2015, p. 19).

An additional factor that influences supply is the physical capacity of the transmission grid, which determines the maximum amount of electricity that can reach consumers. Even when sufficient power is produced, congestion in the high-voltage network can limit transfers between regions, requiring TSOs to intervene through redispatch measures or even grid shutdowns to maintain system stability and prevent overloads (Dannecker, 2015, p. 23). In such situations, the physical limits of the grid – or transmission congestion – lead to regional price differences, often determined by Locational Marginal Pricing (LMP), as constrained areas experience price variations depending on local supply and demand conditions. LMP is a pricing method where the price of electricity is calculated at different points across the transmission grid, instead of using one single price for different regions (Edwards, 2023, p. 133; Müsgens & Bade, 2024, p. 107). As this study focuses on price forecasting within a single market, the detailed modelling of cross-border dynamics will not be examined.

Given these structural, technical, and market-driven dynamics, the functioning of electricity markets is increasingly shaped by the regulatory frameworks that govern investment incentives, system reliability, and market behaviour. Policy instruments are essential to stabilize operations and maintain security of supply amid growing renewable penetration and heightened price volatility (IEA, 2025, pp. 7–9). Regulators have expanded the use of mechanisms such as capacity remuneration schemes to pay generators for their available capacity, ensuring sufficient dispatchable power remains available as variable renewables grow (IEA, 2025, p. 84). Simultaneously, authorities are tightening transparency and monitoring obligations – illustrated by the EU's REMIT framework – to mitigate risks linked to market manipulation and algorithmic bidding, thereby preserving market integrity in short-term markets (IEA, 2025, p. 29; Müsgens & Bade, 2024, p. 173). Current policy discussions emphasize the need for market designs and regulatory tools to align, ensuring affordability, investment certainty, and long-term system reliability (IEA, 2025, p. 15).

Beyond domestic regulation, geopolitical factors and international energy policies play an important role in shaping electricity market conditions. Decisions by major energy-producing countries and organisations, such as OPEC, influence global oil and gas prices, which indirectly affect electricity generation costs in fuel-dependent systems through price transmission mechanisms (Mack, 2014, p. 182; Edwards, 2023, p. 162). In European power systems, fossil fuel prices are particularly relevant because thermal generation technologies are often positioned at the margin of the merit order and therefore contribute directly to short-run price formation when renewable output is insufficient (Edwards, 2023, pp. 28–29). Geopolitical developments, including supply disruptions, sanctions, or changes in export policy, can thus trigger abrupt revisions in fuel price expectations and amplify price volatility in electricity markets, even though their effects materialise indirectly through input costs rather than through electricity demand itself (IEA, 2025, p. 7; Mack, 2014, p. xx).

Taken together, electricity price formation reflects the interaction of demand patterns, generation costs, network constraints, regulatory frameworks and external macroeconomic and geopolitical influences. These factors do not operate in isolation but interact in a dynamic and sometimes non-linear manner, contributing to the pronounced volatility observed in electricity markets. Recognising this complex risk environment is essential for analysing electricity prices and provides the foundation for the forecasting approaches examined in the subsequent sections.

2.1.5 Motivation and common approaches to power trading

As demonstrated by Müsgens and Bade (2024) and Edwards (2023) electricity market instruments exhibit strong parallels to those used in traditional financial markets. While the focus of this study is on spot and day-ahead data, the market structure is also comprised of products such as forwards, futures, and other derivatives. Consequently, this section will highlight these other products to provide a comprehensive introduction but will focus primarily on the short-term spot markets.

As established in Section 2.1.3, the day-ahead and intraday markets form the core of short-term electricity trading. The current section analyses the trading behaviours and motivations of market participants within these venues. The inherent properties of electricity – its non-storability, high volatility, and sensitivity to real-time information –

are the fundamental drivers shaping the distinct and specialized trading strategies employed by market participants (Edwards, 2023, pp. 49–50; Müsgens & Bade, 2024, p. 23).

Motivations differ substantially across actor groups. Utilities and large industrial consumers primarily use short-term markets for hedging, securing expected consumption at predictable prices and thereby reducing exposure to fluctuations in spot prices and standard load profiles, which are "the average (quarter-hourly) power consumption of a typical household" (Müsgens & Bade, 2024, pp. 39–40). To complement these short-term hedges, many large consumers enter into forward contracts, which allow them to stabilise procurement costs over longer horizons by fixing part of their expected load at predetermined prices (IEA, 2025, pp. 62–63). Generation companies align their trading activity with operational optimisation. They monitor indicators such as clean spark and dark spreads, which compare electricity prices with fuel and CO_2 costs, to assess whether operating a power plant yields a positive margin. These spreads reflect the basic economics of thermal generation, where operators pay for fuel and CO_2 certificates while receiving revenue from electricity sales, making the spread a direct measure of short-term profitability. These signals inform dispatch decisions and influence whether conventional units are ramped up, reduced or temporarily shut down. Because these spreads replicate the marginal profit of specific generation technologies, they are also traded OTC as financial products, allowing market participants to hedge or speculate on relative movements between power, fuel and emission prices (Müsgens & Bade, 2024, pp. 66–67; Edwards, 2023, p. 79).

In addition to spot and forward markets, many actors rely on exchange-traded futures and options, which offer liquidity, transparency and central clearing. Utilities and retailers predominantly use futures to manage portfolio risk in a more flexible manner than with bilateral forwards, as futures allow frequent rebalancing through daily margining and can be unwound at low transaction costs in liquid trading sessions (Mack, 2014, pp. 133–134; IEA, 2025, pp. 62–63). Generators employ futures to hedge aggregated production volumes over broader periods rather than individual delivery windows, benefiting from standardised contract sizes that simplify portfolio hedging across multiple months or seasons (Mack, 2014, pp. 141–142). Options are used when actors seek asymmetric protection: market participants buy call or put options to insure against extreme price movements, while option writers earn premiums when expecting stable market conditions

(Mack, 2014, p. 143). Derivative markets therefore enable both hedgers and speculators to manage or assume risk efficiently while contributing to market liquidity (Mack, 2014, p. 167).

Other speculators in electricity markets include proprietary traders, who engage solely for financial returns without any link to physical delivery (Müsgens & Bade, 2024, p. 35), as well as power plant operators, who may pursue speculative opportunities in addition to their standard activity of producing and selling electricity. Proprietary traders and speculative desks adopt opportunity-driven strategies that respond to short-term market fluctuations. Such fluctuations often arise from unexpected forecast revisions or rapid changes in renewable production, which can create temporary price distortions in intraday spot markets (Müsgens & Bade, 2024, pp. 88–90). Their objective is to generate profit from the inherent price volatility in electricity markets, where even slight discrepancies between real-time generation and consumption can prompt price fluctuations, which in turn can be used as trading opportunities (Edwards, 2023, pp. 51 & 241).

The growing operational complexity of electricity systems has also encouraged the use of algorithmic order execution within the intraday market. Automated systems monitor updated weather forecasts, TSO data and plant availability, enabling rapid adjustments to intraday positions and helping to minimise balancing costs for portfolio managers. Such algorithms are often used to close deviations in balancing groups immediately following forecast changes. Nevertheless, discretionary trader judgement remains essential when market conditions deviate from typical patterns or when qualitative information, such as unexpected outage announcements or operational restrictions, must be interpreted (Müsgens & Bade, 2024, pp. 153–154).

Given these behavioural and technological aspects, trading outcomes are further shaped by the risk environment characteristic of electricity markets, as outlined in Section 2.1.4. High intraday volatility, frequent price spikes and uneven liquidity across products require disciplined management of mark-to-market exposures and position limits (Mack, 2014, p. 8). As Mack (2014, p. 8) notes, traders must account not only for price risk but also for operational uncertainties arising from outages and transmission congestion, as these events can rapidly alter expected market conditions. The need to react to such risks in real time places considerable emphasis on timely information processing and on the ability to adjust trading strategies swiftly in response to changing fundamentals.

Taken together, these trading practices show that electricity markets depend heavily on rapid information processing, continuous position management and flexible optimisation. They also provide a strong motivation for incorporating news-based and machine-learning approaches into price forecasting, as examined in the following sections.

2.2 Fundamentals of Natural Language Processing

This section provides a methodological foundation for the construction of feature-based news content. The text introduces the role of NLP in analysing unstructured text (Section 2.2.1), the transformation of text into quantitative embeddings (Section 2.2.2), and the specific challenges of applying these methods within energy market contexts (Section 2.2.3).

2.2.1 Introduction to NLP

NLP is a subfield of machine learning that studies how computers can be used to analyse, interpret and generate human language in textual or spoken form (Lee, 2023, pp. 3–4). In contrast to structured numerical data, natural language is inherently unstructured, ambiguous and context dependent. The meaning of linguistic expressions depends not only on individual words but also on grammatical structure, surrounding context and domain-specific conventions, which makes automated processing particularly challenging (Lee, 2023, p. 5).

In economic and financial research, NLP has gained importance due to the increasing availability of large volumes of textual information, including news articles, policy statements and corporate disclosures, which contain information that is not immediately reflected in quantitative market data (Das, 2014, pp. 3–4). Financial markets are influenced not only by realised economic outcomes but also by expectations, uncertainty and narratives communicated through language. NLP provides a systematised framework for the conversion of unstructured textual information into quantitative representations. These representations can then be integrated into empirical models and decision-making processes, which can be employed for the generation of trading signals (Das, 2014, pp. 4–5).

Earlier approaches relied primarily on frequency-based representations, whereas modern NLP increasingly applies neural language models that can capture contextual meaning and semantic relationships more effectively (Jurafsky & Martin, 2025, p. 4). Recent progress in representation learning, particularly through transformer-based architectures, has substantially improved the ability of NLP models to account for context-dependent semantics (Jurafsky & Martin, 2025, p. 171). These models produce dense numerical representations, commonly referred to as embeddings, which encode both lexical meaning and contextual relationships within a text. Such advances have expanded the applicability of NLP to information-sensitive domains such as finance and energy markets, where small changes in wording may signal shifts in expectations or market conditions (Lee, 2023, p. 118; Das, 2014, p. 4; Simonian, 2025, p. 127).

2.2.2 Conceptual Foundations of Text-to-Quantitative Transformation

A fundamental task in Natural Language Processing is the transformation of unstructured textual data into numerical representations that can be analysed using statistical and machine-learning methods. Unlike numerical variables, text does not possess an inherent quantitative structure, which requires explicit encoding schemes that preserve linguistic meaning while enabling mathematical processing (Jurafsky & Martin, 2025, p. 4). The objective of text representation is therefore to map words, sentences or documents into vector spaces in which semantic similarity can be expressed through numerical distance (Jurafsky & Martin, 2025, p. 99).

Early approaches to text representation rely on frequency-based techniques, such as bag-of-words and term-frequency measures, which describe documents by counting word occurrences without considering word order or context (Jurafsky & Martin, 2025, pp. 4 & 235). To improve their descriptive power, term weighting schemes such as Term Frequency–Inverse Document Frequency (TF-IDF) assign higher importance to words that are frequent within a document but rare across the corpus, thereby reducing the influence of common yet uninformative terms (Jurafsky & Martin, 2025, pp. 235–236; Simonian, 2025, p. 129; Das, 2014, pp. 13 & 81). While these methods are transparent and computationally simple, they fail to capture semantic relationships between words and texts. More advanced representation methods address this limitation by embedding textual elements into continuous vector spaces based on their contextual usage, thereby

allowing semantically related expressions to be represented more closely (Jurafsky & Martin, 2025, p. 95; Lee, 2023, p. 118). Embedding-based representations therefore outperform frequency-based approaches because they capture meaning through contextual co-occurrence patterns rather than isolated word counts. Words or texts that appear in similar contexts are mapped to nearby positions in the vector space, even if they do not share identical vocabulary (Jurafsky & Martin, 2025, pp. 4 & 101). In practice, semantic similarity is encoded through geometric relationships, where cosine similarity or Euclidean distance between vectors provides a quantitative measure of conceptual relatedness, enabling downstream tasks such as classification, clustering and forecasting. While Euclidean distance calculates the absolute spatial magnitude between two points, cosine similarity evaluates the angular orientation of vectors, offering a robust metric that remains independent of document length or vector magnitude (Lee, 2023, pp. 326–327).

While embedding-based representations capture semantic similarity through vector geometry, early embedding models assign a single representation to each word, regardless of its linguistic context. Transformer-based language models extend this framework by producing contextual embeddings, where word representations are dynamically adjusted based on surrounding tokens within a sentence (Jurafsky & Martin, 2025, p. 171). This architecture enables the model to learn complex pairwise relationships and long-range dependencies, producing contextual embeddings that adjust dynamically to a word's usage. These dynamic representations encode rich semantic and syntactic information that can be directly applied to classification and topic modelling. By addressing challenges such as polysemy, where a single word can have multiple meanings, transformers are particularly effective in financial applications. They allow for the detection of subtle contextual cues – such as shifts in tone, uncertainty, or narrative framing – that often precede observable market movements (Jurafsky & Martin, 2025, pp. 172 & 206; Das, 2014, p. 4).

While text representation converts language into numerical form, the analysis of this data typically follows two distinct paradigms: classification and topic modelling. Classification is a supervised learning task where documents are assigned to predefined categories based on labelled training data. Typical applications include sentiment classification, event detection or relevance filtering, where the goal is to map textual inputs to known categories (Lee, 2023, pp. 336–337; Jurafsky & Martin, 2025, pp. 61–62). In contrast, topic modelling is an unsupervised approach that identifies latent

thematic structures within a corpus without prior labelling. Rather than assigning texts to explicit categories, topic models identify clusters of semantically related terms that co-occur frequently across documents, thereby revealing dominant themes present in the data (Lee, 2023, p. 125; Das, 2014, p. 42). In financial research, these techniques enable the systematic transformation of qualitative narratives – such as news flows and policy communications – into structured signals that can be linked to quantitative outcomes, making them particularly valuable for discovering emerging risks that have not yet been captured by traditional market indices (Loughran & McDonald, 2020, p. 2; Simonian, 2025, p. 127).

The high dimensionality of modern text embeddings introduces computational and statistical challenges for both classification and topic modelling. Contextual embeddings generated by transformer models typically inhabit high-dimensional vector spaces, which can impede the efficiency and accuracy of clustering algorithms due to the "curse of dimensionality" (Grootendorst, 2022, p. 2). To address this, dimensionality reduction techniques are employed to project these embeddings into a lower-dimensional space while preserving their essential topological structure. The Uniform Manifold Approximation and Projection (UMAP) algorithm is frequently preferred for this task, as it effectively maintains both local and global semantic relationships, ensuring that semantically similar documents remain grouped together in the reduced space (McInnes et al., 2020, pp. 1–2).

Once embeddings are represented in a reduced space, clustering algorithms can be applied to identify groups of semantically related texts. In contrast to supervised classification, clustering aims to uncover latent structure without predefined labels, which is especially relevant when analysing large and heterogeneous document collections (Das, 2014, p. 42). Density-based approaches such as HDBSCAN are well suited for this purpose, as they can identify clusters of varying size and density while explicitly modelling noise, a common feature of real-world textual datasets (Grootendorst, 2022, p. 3; McInnes & Healy, 2017, pp. 1–2). Together, dimensionality reduction and clustering form a critical conceptual bridge between high-dimensional embeddings and interpretable topic structures in NLP applications.

The final step in the transformation process is the extraction of interpretable labels from the identified clusters. While standard TF-IDF evaluates word importance at the

document level, the class-based TF-IDF (c-TF-IDF) procedure treats all documents within a cluster as a single class. This allows the model to calculate a score that reflects how representative a term is of a specific topic by comparing its frequency within the cluster to its frequency across the entire corpus. By identifying words that are unique to specific clusters, c-TF-IDF provides a precise linguistic description of the underlying market narratives or economic themes captured by the model (Grootendorst, 2022, pp. 3–4).

The conceptual process described in this section is implemented using Bidirectional Encoder Representations from Transformers (BERT), a transformer-based language model that jointly performs tokenisation, contextual embedding generation and semantic representation (Devlin et al., 2019, pp. 1–2; Jurafsky & Martin, 2025, p. 171). BERT encodes linguistic meaning as dense vectors that capture both semantic and syntactic relationships through contextualised representations rather than fixed word-level statistics (Jurafsky & Martin, 2025, p. 197).

This research paper employs specialised variants of the BERT architecture that are tailored to the computational and analytical requirements of large-scale text analysis. Compact models such as MiniLM reduce computational complexity while preserving much of the semantic representation capacity, making them suitable for large-scale text analysis (Wang et al., 2020, p. 1). In addition, zero-shot classification models based on transformer embeddings enable semantic categorisation without task-specific labelled data, relying on general language understanding acquired during pretraining (Yin et al., 2019, pp. 1–2). The specific implementation choices and empirical integration of MiniLM and zero-shot are presented in detail in Chapter 3.

2.2.3 Limitations and Challenges of NLP in Financial Applications

Despite its growing relevance, the application of Natural Language Processing in financial contexts faces several methodological and practical limitations. A central challenge arises from the ambiguity and context sensitivity of financial language, where identical terms may carry different meanings depending on macroeconomic conditions, institutional settings or market regimes. This linguistic complexity increases the risk of misinterpretation, particularly when models are applied across heterogeneous sources

such as news articles, policy statements and analyst commentary (Das, 2014, pp. 105–107).

Another limitation concerns data quality and representativeness. Financial text data often contain noise, redundancy and reporting bias, as media coverage tends to concentrate on extreme events or prominent institutions. Such distortions can lead NLP models to overemphasise sensational narratives while underrepresenting gradual or structural developments, potentially weakening the stability of derived signals (Loughran & McDonald, 2020, p. 3; Das, 2014, p. 105). Furthermore, changes in language use over time, driven by regulatory shifts or evolving market discourse, may reduce model performance if not adequately addressed through retraining or adaptive techniques (Loughran & McDonald, 2020, pp. 4–5).

From a modelling perspective, transformer-based architectures introduce challenges related to interpretability and computational cost. While contextual embeddings offer superior semantic representation, their internal decision processes remain largely opaque, complicating economic interpretation and model validation (Wang et al., 2020, p. 1; Das, 2014, p. 107; Jurafsky & Martin, 2025, p. 172). In addition, the reliance on large pretrained models raises concerns regarding domain mismatch, as models trained on general language corpora may not fully capture specialised financial or energy-market terminology without further adaptation (Kaplan et al., 2023, p. 1; Das, 2014, p. 105; Jurafsky & Martin, 2025, p. 162).

Taken together, these limitations highlight the importance of careful model selection, robust preprocessing and complementary quantitative validation when deploying NLP techniques in financial and energy market research.

2.3 Fundamentals of Trade Signal Generation using NLP

Building upon the linguistic foundations of Section 2.2, this section examines the practical application of NLP within financial and energy markets. The text distinguishes between two concepts: statistical forecasting (Section 2.3.1) and actionable trading implementation (Section 2.3.2). The former focuses on demonstrating the predictive power of news beyond quantitative inputs, while the latter addresses the transformation of these forecasts into directional signals using threshold rules and risk controls. Attention

is given to empirical applications in energy and electricity markets, where rapid information processing and expectation-driven price dynamics make NLP-based trading strategies especially relevant.

2.3.1 From Textual Information to Market Forecasts

Building on the power market characteristics discussed in Section 2.1 and the text-to-vector transformation outlined in Section 2.2, recent research integrates textual information into forecasting models for energy and commodity prices. Standard forecasting setups typically rely on historical prices, load patterns, weather variables, and physical system constraints, which capture structural and seasonal behaviour but only partially reflect forward-looking information (Hong et al., 2020, pp. 377–378; Müsgens & Bade, 2024, pp. 5–6). Textual sources such as news and policy communication can contain expectation- and uncertainty-related signals that precede observable price adjustments, making them a natural complement to purely quantitative inputs (Das, 2014, p. 4; Loughran & McDonald, 2020, pp. 1–2). In energy markets, this channel is particularly relevant because regulatory and geopolitical developments may shift price expectations before physical effects materialise (Dannecker, 2015, pp. 2–3; Edwards, 2023, p. 162).

Early empirical studies linking news to market outcomes focused on sentiment indicators derived from textual data. Lee Smales (2014) examines the gold futures market using commodity-specific news sentiment scores obtained from structured news analytics. The study aggregates sentiment indicators extracted from financial media and incorporates them into regression-based forecasting models alongside historical price information news (Smales, 2014, pp. 277–278). Using daily gold futures data over an extended sample period, Smales (2014, pp. 285–286) shows that news sentiment significantly improves the prediction of short-term returns and volatility, with negative news exerting a stronger market impact than positive news.

Similarly, Li et al. (2021) analyse oil futures markets and construct sentiment measures from oil-related news articles. Their framework combines textual sentiment indicators with traditional econometric volatility models, demonstrating that sentiment contains incremental predictive information for both returns and volatility, particularly during periods of elevated uncertainty (Li et al., 2021, pp. 3–4). These early contributions treat

sentiment as a low-dimensional numerical signal derived from text, consistent with the frequency-based and term-weighting approaches discussed in Section 2.2.

Subsequent research extends this idea by embedding textual features directly into forecasting models rather than relying solely on aggregated sentiment scores. Guan et al. (2022) propose a forecasting framework for international natural gas prices that combines historical price series with numerical representations of related news articles. Textual information is first transformed into quantitative features and then integrated into machine-learning models alongside lagged price inputs (Guan et al., 2022, pp. 5–6). Their empirical results show that news-augmented models outperform price-only benchmarks, especially in capturing abrupt price movements associated with geopolitical events and supply disruptions (Guan et al., 2022, p. 10). This approach illustrates how text embeddings, as introduced in Section 2.2, can be treated as additional explanatory variables within standard forecasting pipelines.

More recent studies leverage advances in transformer-based language models to extract richer semantic representations from news. Chakraborty et al. (2024, pp. 3–4) develop an event-driven forecasting framework in which news articles are encoded using contextual embeddings derived from transformer models. Their Robust Event Network (REN) model is a neural network-based predictive model that specifically utilizes Long Short-Term Memory (LSTM) hidden states to capture non-linear, long-term dependencies between extracted event intensities and commodity price fluctuations (Chakraborty et al., 2024, pp. 6 & 13). Rather than modelling sentiment alone, this framework explicitly links semantic content to price reactions. This aligns with the argument of Kaplan et al. (2023) that general language models require asset-specific adaptation to capture economically relevant information in energy and commodity markets (Kaplan et al., 2023, p. 1). The authors demonstrate that incorporating exogenous event features from news significantly improves accuracy for commodities sensitive to real-world developments. They highlight that while traditional models often fail to account for extreme volatility, their approach captures the impact of external shocks and provides early warning signals for abrupt price fluctuations following significant information releases (Chakraborty et al., 2024, pp. 24–25).

While RENs rely on deep learning architectures, other high-performing frameworks in financial forecasting frequently employ gradient-boosted decision trees to integrate these

textual embeddings with numerical data. Methods such as XGBoost are particularly effective in this context as they iteratively combine multiple weak learners into a strong predictive model through additive training, minimizing forecast errors while efficiently capturing complex, non-linear interactions (T. Chen & Guestrin, 2016, pp. 1–2; Simonian, 2025, pp. 54 & 65).

Within electricity markets, the relevance of NLP-based forecasting has been demonstrated more explicitly. Bai et al. (2024, pp. 1–2) investigate the joint role of news and load information in electricity price forecasting. Their framework combines traditional physical variables, such as demand and load forecasts, with textual features extracted from energy-related news. Using supervised learning models, they show that incorporating news-derived signals improves forecast accuracy relative to models relying solely on load, weather and historical prices. The results underline that news captures regulatory, infrastructural and geopolitical information that is not directly observable in physical system data, reinforcing the complementarity between NLP-based features and established electricity market fundamentals (Bai et al., 2024, pp. 5–6).

Bi et al. (2025, pp. 1–2) further advance this line of research by introducing a large-scale, news-augmented benchmark dataset for electricity price forecasting using large language models. News texts are encoded using transformer-based embeddings and integrated into forecasting architectures across multiple time horizons (Bi et al., 2025, p. 6). Their empirical analysis demonstrates that NLP-augmented models deliver improvements in predictive performance, particularly during periods of market stress and heightened volatility (Bi et al., 2025, pp. 7 & 9). This study provides systematic evidence that textual representations derived from large language models are not merely auxiliary inputs but can become central components of modern electricity forecasting frameworks (Bi et al., 2025, pp. 2–3).

Taken together, the literature demonstrates a clear evolution from simple sentiment indicators towards embedding-based and event-driven forecasting frameworks. Textual information is increasingly treated as a quantitative input that complements historical prices and physical fundamentals. In energy markets in particular, where prices react strongly to expectations and policy-driven events, NLP-based forecasting models provide a natural extension to established time-series approaches discussed by Dannecker (2015)

and Müsgens and Bade (2024). This body of research forms the conceptual foundation for the signal-generation framework developed in the subsequent Chapter 3.

Maybe add a paragraph about the limitations?

2.3.2 From Forecasts to Trading Signals

While the literature reviewed in Section 2.3.1 focuses on improving predictive accuracy, forecasting performance alone does not automatically translate into economically meaningful trading outcomes. The transformation of forecasts into trading signals requires an additional decision layer that accounts for market frictions, risk constraints, and asymmetric payoff structures. In practice, market participants must determine when a forecast is sufficiently strong to justify a position and how it should be implemented within the specific microstructure of the traded market (Mack, 2014, pp. 7–9; Müsgens & Bade, 2024, pp. 22–23).

A common approach in the literature is to convert point forecasts or directional predictions into discrete trading actions, such as long, short, or neutral positions. This step typically involves the use of confidence thresholds or probability cut-offs, whereby positions are only entered when the model's signal exceeds a predefined level of certainty. Such threshold-based decision rules help mitigate overtrading and reduce exposure to noise-driven forecasts, a consideration that is particularly important in volatile and illiquid segments of electricity markets (Müsgens & Bade, 2024, pp. 33 & 44; Chakraborty et al., 2024, p. 7; Edwards, 2023, pp. 51 & 270).

A closely related application of this forecast-to-signal paradigm can be found in electricity markets. Chen et al. (2025) analyse intraday electricity price forecasting from an explicitly economic perspective by embedding probabilistic price forecasts into concrete trading strategies, rather than evaluating models solely on statistical accuracy. Specifically, the authors translate forecasted price distributions into execution rules such as optimal selling strategies under fixed-volume and time constraints, thereby explicitly modelling the link between forecast uncertainty and trading decisions (J. Chen et al., 2025, pp. 1–3). Their empirical results show that improvements in forecast quality do not automatically lead to higher trading profits, as realised performance depends critically on execution timing, liquidity conditions, and market constraints in the intraday setting (J. Chen et al., 2025, pp. 18–19). These findings highlight that, even in electricity markets with rich physical fundamentals, the transformation of forecasts into tradable signals

constitutes a distinct optimisation problem that must be assessed using economic performance criteria rather than predictive accuracy alone (J. Chen et al., 2025, p. 21).

Extending this forecast-to-execution separation beyond electricity markets, Kirtac and Germano (2024, pp. 3–4) apply a similar two-stage framework to the stock market by first producing sentiment-based directional predictions using transformer embeddings and subsequently applying confidence filters to derive trading positions. Methodologically, the authors derive sentiment scores from large language models that are fine-tuned using future excess returns as supervision, thereby directly linking textual information to expected price direction. Trading strategies are then constructed as long–short portfolios based on percentile thresholds of these sentiment scores, with transaction costs explicitly incorporated to assess economic viability under realistic market conditions. Their empirical results show that profitability is highly sensitive to signal calibration. Strategies with stricter confidence thresholds exhibit lower trade frequency but achieve superior risk-adjusted performance, as measured by Sharpe ratios and drawdown metrics (Kirtac & Germano, 2024, pp. 7–8).

In contrast to stock markets, the conversion from forecasts to trades in energy markets must also reflect physical and institutional constraints. Electricity prices are subject to sharp intraday reversals, negative pricing, and uneven liquidity across delivery periods (Section 2.1). As noted by Müsgens and Bade (2024, pp. 89 & 153) and Chakraborty et al. (2024, pp. 1 & 24), trading strategies that ignore these features risk systematic losses despite statistically accurate forecasts. Consequently, practical implementations often restrict trading to specific time windows, delivery products, or volatility regimes, aligning model outputs with operational feasibility. Interpretability plays a particularly important role in this context. Electricity traders frequently rely on discretionary judgement when reacting to unexpected events such as outages, regulatory announcements, or geopolitical developments. Therefore, qualitative understanding of why a price signal arises is often as important as the signal itself. Topic-based representations derived from c-TF-IDF, as introduced in Section 2.2.2, support this requirement by providing transparent keyword descriptions for each identified topic. This allows traders to distinguish, for example, between signals driven by "supply outages" and those linked to "macroeconomic conditions" (Grootendorst, 2022, pp. 3–4). Such interpretability facilitates discretionary overlays and reduces the risk of mechanically executing opaque model outputs.

A critical practical limitation in electricity markets concerns signal timeliness. As noted in Section 2.2.3, the substantial computational costs associated with transformer-based models can introduce latency between news arrival and signal generation. In the intraday market, where price adjustments may occur within minutes of an information release, this creates a trade-off between semantic richness and signal velocity (Müsgens & Bade, 2024, p. 153; Edwards, 2023, p. 51). More complex models risk producing 'stale' signals that decay before execution becomes economically viable, necessitating a balance between representational depth and computational efficiency in short-horizon contexts (Das, 2014, pp. 106–107; Jurafsky & Martin, 2025, p. 172).

Taken together, the literature demonstrates that the economic value of NLP-based forecasts depends critically on the design of the signal-generation layer rather than on predictive accuracy alone. Across both electricity and stock markets, effective strategies do not translate forecasts into trades mechanically but embed them within disciplined execution frameworks that account for confidence calibration, market microstructure, interpretability, and latency constraints. Empirical evidence shows that forecast improvements only become economically meaningful when combined with selective execution rules, risk controls, and market-specific feasibility considerations. This distinction between statistical forecasting performance and tradable signal construction motivates the two-stage modelling approach adopted in this study. Consistent with the reviewed literature, NLP-derived forecasts are first generated and subsequently transformed into actionable long, short, or hold decisions under explicit decision rules, as detailed in Chapter 3.

3 Methodology

XXX

3.1 Framework Structure

- Introduction to our Data & Methodology pipeline
- Showcase an overview of the process used (flowchart) from data sourcing all the way to model evaluation & backtesting.
- Explain and highlight which models will be used (XGBoost, Ridge & Logistic Regression). We use a dual model set-up, where we first classify the next day spread to be positive or negative alongside a confidence score, then use a Logistic (or lightGBM) to then turn those forecast into long/short/hold decisions
- Quickly explain how the infrastructure was set up (GPU clustering, code optimisation for parallel processing)
- Give our model a name (ex. PowerTree)

3.2 Data Sourcing

This research project combines two synchronised data streams: (i) news headlines capturing energy-related information flows and (ii) German electricity market data capturing realised price and volume dynamics. The sourcing design prioritises coverage, timestamp quality, and hourly alignment, enabling subsequent aggregation of news signals to the market sampling frequency (SOURCE). Sections 3.2.1 and 3.2.2 describe the collection of news and market data, respectively.

3.2.1 News Data Collection

News data were collected using NewsAPI.org, an aggregation service providing programmatic access to historical and current news articles from diverse sources. NewsAPI.org was selected over alternative services due to its comprehensive metadata structure, including timestamps, source attribution, and article descriptions, which enables precise temporal alignment with market data. The service's structured data format

has been employed in recent commodity market forecasting applications to extract market-relevant signals from high-frequency news flows (SOURCE).

The news collection strategy encompassed both English- and German-language sources to capture the full spectrum of discourse surrounding German electricity markets. German-language sources provide domestic market insights, regulatory announcements, and regional policy developments, whilst English-language international sources offer broader market analysis, global energy trends, and institutional investor perspectives that influence local prices (SOURCE). This bilingual approach is particularly relevant given Germany's central role in European electricity markets, where cross-border flows and interconnector capacity create price dependencies across multiple national markets (SOURCE). All collected articles were subsequently translated into German (as detailed in Section 3.3.1) to enable consistent downstream natural language processing.

The collection period spanned from 1 November 2020 to 31 October 2025, providing a five-year window that captures significant market regime changes: the COVID-19 pandemic's demand disruptions, the 2022–2023 energy crisis triggered by geopolitical tensions following Russia's invasion of Ukraine, and the subsequent market stabilisation accompanied by accelerated renewable deployment (SOURCE). This temporal coverage exposes the model to diverse market conditions, including both low-volatility and crisis periods, thereby improving robustness to the regime changes that characterise modern electricity markets.

The retrieval strategy employed daily queries via NewsAPI.org's premium subscription tier, which permits up to 100 articles per API call. Each day's query retrieved articles matching the specified keywords and source parameters (as detailed in Table 1). On days when more than 100 matching articles were published, the API returned the 100 most relevant results ranked by NewsAPI's relevance algorithm; however, such instances were rare over the five-year period. This systematic daily collection methodology aligns with best practices for constructing news-based trading signals, where consistent temporal coverage prevents information gaps that could compromise signal quality (Smales, 2014, pp. 277–279). The collection contained 89,517 English-language and 122,208 German-language articles before any preprocessing.

Table 1: News Search Parameters

Parameter	*English Sources*	*German Sources*

News Outlets	ABC News, Associated Press, BBC News, Bloomberg, CNN, MSNBC, NBC News, Reuters, The Wall Street Journal, The Washington Post, Time, USA Today	Bild, Der Tagesspiegel, Die Zeit, Focus, Handelsblatt, Spiegel Online, WirtschaftsWoche
Keywords	EPEX, Erdgas, Erdöl, EU, EZB, IEA, Kohle, OPEC, Strombedarf, Stromeffizienz, Stromerzeugung, Stromnachfrage, Stromsparen, Wetter, Zentralbank	ACER, Carbon, Climate, Coal, ECB, Electricity, Emissions, Energy, EPEX, European Comission, European Union, Gas, Grid, Hydro, IEA, Infrastructure, Interest Rates, Market, Nuclear, OECb, OPEC, Oil, Pipeline, Policy, Renewable, Solar, Storage, Supply, Tariffs, Transmission, Wind
Articles Retrieved	89’517	122’208
Timeframe	1 November 2020 - 31 October 2025	1 November 2020 - 31 October 2025

The keyword selection strategy balanced domain specificity with comprehensive coverage, incorporating terminology from energy infrastructure, market mechanisms, relevant institution abbreviations, and macroeconomic indicators. This multi-faceted approach recognises that electricity prices respond to diverse information sources, ranging from technical factors such as renewable generation forecasts to broader economic indicators that affect industrial demand (QUELLE AUS TEIL OBEN).

3.2.2 Energy Market Data Collection

Energy market data were sourced from Energy-Charts.info, a platform maintained by the Fraunhofer Institute for Solar Energy Systems (ISE) that has become a widely used resource in German electricity market research (Wirth, 2023). The platform aggregates data from multiple authoritative sources, including ENTSO-E, EPEX SPOT, and Germany's four TSO’s, providing harmonised datasets that would otherwise require complex reconciliation across different source systems (Fraunhofer ISE, 2024; Schmidt et al., 2023).

Data were retrieved at 15-minute granularity for the period 1 November 2020 to 31 October 2025, matching the news collection timeframe. The following price variables were collected: day-ahead auction prices (EUR/MWh) representing the uniform clearing prices established in the daily EPEX SPOT auction, and intraday continuous market prices (EUR/MWh) reflecting real-time trading activity. These two price series form the basis for computing the spot-day-ahead spread that serves as the target variable for this study.

Power generation volumes were retrieved, disaggregated by source category to capture the supply-side dynamics that influence price formation. The dataset includes output from onshore wind, offshore wind, solar photovoltaic, nuclear, lignite, hard coal, natural gas, pumped-hydro, and run-of-river hydro facilities (SOURCE). This granular breakdown enables the model to account for the merit-order effects discussed in Section 2.1, where low-marginal-cost renewables displace higher-cost thermal generation and compress electricity prices during periods of high renewable feed-in. Total electricity demand (MW) was also collected, representing the aggregate load that must be balanced against generation in real time.

The raw dataset comprised approximately 175,000 observations at 15-minute resolution. These data were converted from the original JSON format to CSV for subsequent preprocessing. Quality assessment of the raw data identified several categories of missing or anomalous values. Gaps in price reporting occurred during approximately 0.4% of 15-minute intervals, concentrated around daylight saving time transitions, exchange system maintenance windows, and periods of extreme market stress when trading was suspended (SOURCE).

3.3 Data Preprocessing

The preprocessing pipeline comprised two parallel workflows: one for news data and one for energy market data, which were subsequently integrated to construct the final analytical dataset. This section details the procedures employed to ensure data quality, topical relevance, and temporal alignment across these heterogeneous sources.

3.3.1 News Corpus Refinement

The broad keyword-based retrieval strategy employed in Section 3.2.1, whilst ensuring comprehensive coverage of potentially relevant content, inevitably captured substantial quantities of off-topic material. Articles matching terms such as "energy" or "market" frequently addressed topics unrelated to electricity trading, including sports events, entertainment news, and general political coverage. Rigorous filtering was therefore required to isolate energy-relevant news content suitable for signal generation.

Off-topic content was identified and removed through an iterative refinement procedure using BERTopic (Grootendorst, 2022, pp. 1–3), an unsupervised topic modelling

framework described in Sectio 2.2.3. BERTopic was selected for this task because it enables the discovery of latent thematic structures without requiring predefined category labels, making it well-suited to exploratory corpus analysis in which the full range of off-topic content is not known in advance (SOURCE). The framework's use of HDBSCAN clustering also accommodates topics of varying sizes and explicitly identifies noise documents that do not belong to any coherent theme.

Each iteration of the refinement procedure proceeded as follows: (i) a topic model was fitted to the corpus to discover latent thematic clusters; (ii) human reviewers inspected topic labels, representative keywords, and sample documents to identify non-relevant themes; (iii) characteristic keywords and phrases associated with off-topic clusters were extracted to construct exclusion filters; (iv) articles containing these keywords in titles or descriptions were removed from the corpus; and (v) the filtered corpus was resubmitted for subsequent refinement. This human-in-the-loop approach ensured that automated clustering was validated against domain expertise before content removal, preventing the erroneous exclusion of legitimately relevant articles that might share surface-level vocabulary with off-topic content.

Nine iterative filtering rounds were applied to the English corpus, progressively removing articles associated with sports coverage (team names, match results, athlete profiles), entertainment news (film releases, celebrity coverage, television programming), crime reporting (arrests, trials, investigations), retail and consumer news (product launches, shopping trends), health and medical coverage unrelated to energy policy, weather reporting without market implications, and general political coverage lacking energy or economic relevance. This process reduced the English corpus from 89,517 to 39,744 articles, a 55.6% reduction. The substantial reduction reflects the broad semantic scope of English-language international news sources, where energy-related keywords frequently co-occur with unrelated content.

The German corpus required only three iterative rounds, targeting crime, education, sports, and tourism content, and was reduced from 122,208 to 113,894 articles, which represents a more modest 6.8% reduction. The disparity in reduction rates suggests that German-language sources retrieved via energy-specific keywords were more topical than their English-language counterparts. This likely reflects the more specialised focus of German business publications such as Handelsblatt and WirtschaftsWoche, which

concentrate on economic and market coverage rather than general news. A detailed log of excluded topic categories and representative keywords per iteration is provided in Appendix X.

Following topic-based filtering, English articles were translated into German using the Google Translate API via the deep-translator Python library. Translation into German (rather than translating German content into English) was chosen for two reasons. First, German-language transformer models fine-tuned on energy and economic terminology prevalent in the DACH-region discourse were available for downstream classification tasks. Second, the majority of the filtered corpus (74%) was already in German, minimising the proportion of content subject to potential translation errors. The translated English corpus (39,744 articles) was merged with the filtered German corpus (113,894 articles), yielding a total of 153,638 articles. After removing translation failures and malformed entries, the final corpus comprised 149,512 German-language articles spanning the five-year collection period.

3.3.2 Energy Market Data Processing

Energy market data underwent preprocessing to ensure temporal alignment with the news corpus and to prepare variables for feature engineering. The raw data, retrieved at 15-minute granularity, were resampled to hourly frequency by computing means for price variables and sums for volume variables. Hourly aggregation was selected to align with the day-ahead auction structure, where prices are determined for each hour of the following day, and to reduce computational overhead whilst preserving meaningful intraday variation. Following resampling, the energy dataset comprised 43,824 hourly observations spanning the five-year collection period.

Missing values were identified in 192 hourly records, primarily during daylight saving time transitions and periods of data reporting gaps. These observations were excluded rather than imputed to preserve data integrity, as interpolation of price data during market discontinuities could introduce spurious patterns. Anomalous price observations, including negative prices below -500 EUR/MWh and favourable prices exceeding 1,000 EUR/MWh, were retained in the dataset, as these extreme values represent genuine market conditions during periods of renewable oversupply or acute scarcity rather than data errors. The final processed energy dataset comprised 43,632 hourly observations.

3.4 Feature Engineering

The feature engineering process transforms raw energy market data and pre-processed news articles into a structured feature matrix suitable for forecasting modelling. This section details the construction of baseline features, the extraction of topical signals via zero-shot classification, the generation of semantic embeddings, and their temporal aggregation via exponential decay weighting.

3.4.1 Baseline Features

The baseline feature set comprises variables derived directly from the energy market data described in Section 3.2.2. These features capture the fundamental supply-demand dynamics that drive electricity price formation, as well as temporal patterns that reflect the cyclical nature of electricity consumption (SOURCE 2.1).

Price and Spread Features: The primary price variables are the day-ahead auction price and the intraday spot price, resampled at hourly frequency. The absolute spread between these prices serves as the target variable for classification. Observations are labelled as Long (positive spread, spot > day-ahead), Short (negative spread, spot < day-ahead), or Neutral when the absolute spread falls within a symmetric deadband of ±3.0 EUR/MWh, which is based on the observed average Bid-Ask spread in the markets (SOURCE).

Generation and Load Features: Aggregate electricity generation and total system load (MW) provide direct measures of supply–demand balance. As discussed in Section 2.1.4, the merit-order effect implies that price formation depends critically on the marginal generating unit dispatched to meet demand (Edwards, 2023, pp. 134–138).

Lagged Features: To capture autocorrelation in electricity prices and demand patterns, lagged values are included at two horizons: 24 hours and 168 hours. Specifically, the feature set contains *price_lag_24*, *price_lag_168*, *load_lag_24*, and *load_lag_168*. These lags reflect the strong daily and weekly seasonality in electricity consumption patterns, as discussed in Section 2.1.4, in which demand exhibits pronounced differences between weekdays and weekends and between daytime peak hours and overnight off-peak periods(Weron, 2014, pp. 1049–1051).

Temporal Features: Categorical and cyclical time indicators encode the position of each observation within relevant periodic structures:

- *hour* (0–23): Captures intraday variation aligned with consumption patterns
- *day_of_week* (0–6): Distinguishes weekday from weekend dynamics
- *day_of_year* (1–366): Encodes annual seasonal effects, including heating and cooling demand
- *week_of_year* (1–52): Provides coarser seasonal resolution
- *month* (1–12): Captures broader seasonal trends in renewable generation and demand

Together, these 13 baseline features establish a foundation that captures the structural characteristics of electricity markets before the incorporation of news-derived signals (SOURCE).

3.4.2 Topic Classification via Zero-Shot Learning

News articles are transformed into structured topical features via a zero-shot classification pipeline that assigns each article to predefined energy-relevant categories derived from Section 2.1. This approach leverages the semantic generalisation capabilities of NLI-based models, as discussed in Section 2.2.2, enabling classification without task-specific training data (Yin et al., 2019, pp. 1–3).

The classification employs the mDeBERTa-v3-base-xnli-multilingual-nli-2mil7 model (Laurer et al., 2024), which extends Microsoft's DeBERTa-v3 architecture (He et al., 2023) to multilingual zero-shot classification. The model outputs probabilities for each label, and the label with the highest probability is assigned as the classification. This enables flexible classification across arbitrary label sets without retraining (Laurer et al., 2024, pp. 3–6). The classification uses the following hypothesis template:

"Der Artikel handelt von: {label}."

The topic structure comprises eleven high-level categories reflecting distinct channels through which news may influence electricity prices. Each category (except "Sonstiges") contains two directionally opposed leaf topics: one representing conditions that would exert upward pressure on prices and one representing conditions that would ease prices, enabling the model to capture both topical relevance and the implied price-impact direction. The complete taxonomy is presented in Table XX.

Category	*Leaf Topics*
Nachfrage	*Rising electricity consumption due to economic activity, industry, or extreme weather; Falling electricity consumption due to economic slowdown or mild weather*
Angebot	*Power plant outages, grid bottlenecks, or low renewable feed-in; High renewable feed-in, new capacity, or stable grids*
Brennstoffpreise	*Rising gas, coal, or CO_2 prices; Falling gas, coal, or CO_2 prices*
Wetter	*Cold, calm, or overcast weather; Mild weather, strong winds, or high solar irradiance*
Wirtschaft & Konjunktur	*Positive economic development, rising industrial production, or corporate growth; Recession, corporate crisis, or declining industrial production*
Finanzmärkte & Geldpolitik	*Interest rate decisions, inflation, or currency fluctuations; Stock market news, corporate earnings, or investments*
Handel & Außenwirtschaft	*Tariffs, trade conflicts, or export restrictions; Trade agreements, market opening, or supply chain developments*
Geopolitik & Konflikte	*War, sanctions, terror threats, or geopolitical tensions; Peace talks, diplomacy, or lifting of sanctions*
Technologie & Industrie	*Technology development, semiconductors, e-mobility, or battery production; Industrial policy, production sites, or corporate strategies*
Politik & Regulierung	*Energy policy, climate legislation, or EU regulation; Domestic politics, government formation, or elections*
Sonstiges	*Sports, entertainment, culture, or local news without economic relevance*

Each article's title is first classified against the eleven category labels. If the highest-scoring category exceeds the Stage 1 threshold ($\tau_1 = 0.25$), the article proceeds to Stage 2 classification within that category; otherwise, it is routed to "Sonstiges." In Stage 2, a threshold of $\tau_2 = 0.20$ governs whether the specific leaf topic is retained. Articles initially classified as "Sonstiges" undergo secondary classification using the article description, providing a fallback for ambiguous headlines.

The threshold values were determined through manual inspection of classification score distributions, balancing coverage against precision(Laurer et al., 2024, pp. 5–7). Articles ultimately classified as having no energy relevance (“Sonstiges”) are excluded from subsequent aggregation.

3.4.3 Sentence Embedding Generation

Complementing discrete topic classification, dense sentence embeddings capture semantic information that predefined categories may not fully represent. As introduced in Section 2.2.2, sentence embeddings encode the full semantic content of text as continuous vectors, preserving nuanced information that discrete labels cannot capture.

The embedding pipeline employs the paraphrase-multilingual-MiniLM-L12-v2 model (Reimers & Gurevych, 2019, pp. 1–2), which builds on the MiniLM knowledge distillation framework (Wang et al., 2020, pp. 2–5) to compress larger transformer models

while preserving most of their representational capacity. This multilingual MiniLM model generates 384-dimensional vector representations optimised for semantic similarity across 50+ languages, including German. For each article, the embedding is computed from the title using mean pooling over contextualised token representations, prioritising computational efficiency while capturing primary informational content.

3.4.4 Time-Decay Aggregation of News Features

Individual article-level features must be aggregated to the hourly resolution of the energy market data. The aggregation method in this research project employs exponential time-decay weighting, assigning greater influence to recent articles while retaining diminishing contributions from older content. This reflects the intuition that news impact decays as information is absorbed into prices (Tetlock, 2007, pp. 1149–1151), which effect has also been observed in the commodity markets (Li et al., 2021, pp. 3–4).

In electricity markets, the high frequency of price dynamics and the rapid incorporation of information suggest that news relevance diminishes relatively quickly. Forecast revisions for renewable generation, updates to plant availability, and shifts in demand expectations are continuously reflected in intraday trading, implying that older news articles provide progressively less marginal information (Weron, 2014, pp. 1031–1032).

The exponential decay formulation provides a parsimonious representation of this information decay:

$$\mathrm{w_i} = \mathrm{e}^{-\lambda \cdot \mathrm{h_i}}$$

Where w_i is the weight assigned to the article i, λ is the decay rate parameter, and h_i is the number of hours elapsed since the article i was published. Higher values of λ produce faster decay, concentrating influence on very recent articles, while lower values extend the effective memory horizon, allowing older articles to retain meaningful weight.

For each hourly timestamp in the energy dataset, time-decayed weighted counts are computed for each of the 14 leaf topics. The aggregation proceeds as follows:

1. **Define Lookback Window**: For a timestamp t, identify all articles published within the interval $(\mathrm{t} - \mathrm{L}, \mathrm{t}]$, where L the lookback window is in hours.

2. **Compute Decay Weights**: For each article i in the window, calculate the decay weight $w_i = e^{-\lambda \cdot h_i}$, where $h_i = t - t_i$ is the number of hours since the time of publication.
3. **Aggregate by Topic**: For each topic k, compute the weighted count as $\text{topic_count}_k(t) = \sum_{i \in \text{topic } k} w_i$
4. **Filter Irrelevant Articles**: Articles classified as "Sonstiges" are excluded from the aggregation, ensuring that topic features reflect only energy-relevant news content.

The output is a matrix of dimension $(T \times 14)$, where T denotes the number of hourly timestamps, and each column represents the time-decayed count for a given topic. This aggregation preserves the interpretability of topical features: higher values indicate greater recent news volume on that topic, with recency weighted. The directional framing of topics (e.g., "gas prices rising" vs. "gas prices falling") enables the model to distinguish between news flows with opposing price implications.

Time-decayed aggregation of embeddings follows a similar procedure, with the key difference that vector averages replace scalar sums:

1. **Define Lookback Window**: Identify articles within $(t - L, t]$ for each timestamp t.
2. **Compute Decay Weights**: Calculate $w_i = e^{-\lambda \cdot h_i}$ for each article.
3. **Weighted Average Embedding**: Compute the aggregated embedding as $e_{agg}(t) = \frac{\sum_i w_i \cdot e_i}{\sum_i w_i}$, where e_i is the 384-dimensional embedding vector for the article i.

The weighted average ensures that the aggregated embedding lies within the convex hull of individual article embeddings, preserving semantic interpretability. If no articles fall within the lookback window for a given timestamp, the aggregated embedding is set to zero, indicating the absence of relevant news content. The output is a matrix of dimension $(T \times 384)$, representing the time-decayed average semantic content of recent news at each timestamp.

The lookback window (L) and decay rate (λ) jointly determine the temporal dynamics of news feature aggregation. Shorter lookback windows with higher decay rates capture

rapid, short-term news effects, while longer windows with lower decay rates smooth over transient fluctuations to capture more persistent trends.

Given the uncertainty regarding optimal parameter values, a grid search approach is employed across a range of parameter combinations:

Parameter	*Values tested*
Lookback Window (L)	24, 48, 72, 168, 336, 504 hours
Decay Lambda (λ)	0.01, 0.05, 0.1, 0.25, 0.5

This yields 30 parameter combinations, each producing a distinct set of time-decayed topic and embedding features. The optimal combination is selected through the model validation procedure described in Section 3.5, which evaluates downstream prediction performance across the parameter grid.

The default parameters ($L = 336$ hours, $\lambda = 0.05$) reflect a two-week lookback with moderate decay, balancing responsiveness to recent news against stability from a broader temporal context. With $\lambda = 0.05$, an article published 24 hours ago retains approximately 30% of its initial weight ($e^{-0.05\times24} \approx 0.30$), while an article from one week ago (168 hours) retains only about 0.02% ($e^{-0.05\times168} \approx 0.0002$).

3.4.5 Dimensionality Reduction of Aggregated Embeddings

The 384-dimensional time-decayed embedding vectors pose challenges for downstream modelling due to high dimensionality and overfitting risk. UMAP, introduced in Section 2.2.2, is employed for dimensionality reduction, offering superior scalability and global structure preservation compared to alternatives (McInnes et al., 2020, pp. 1–2).

The UMAP configuration employs the following parameters:

- ***n_components* = 20**: Target dimensionality, reducing the 384-dimensional input by a factor of approximately 19
- ***n_neighbors* = 15**: Local neighbourhood size, controlling the balance between local and global structure preservation
- ***min_dist* = 0.1**: Minimum distance between points in the embedding, governing clustering tightness

These parameters follow the guidelines provided by McInnes et al. (2018), who recommend *n_neighbors* values between 10 and 50 to balance local and global structure

preservation, and *min_dist* values between 0.0 and 0.5 to control the tightness of the resulting clusters. The choice of 20 components provides sufficient expressiveness for downstream prediction while substantially reducing the dimensionality of the feature space.

WE CANNOT SET RANDOM SEED FOR UMAP FOR PARALLEL PROCESSING!

The final output is a matrix of dimension $(T \times 20)$, where each row represents the reduced embedding for a single hourly timestamp. These 20 embedding dimensions, combined with the 14 topic count features, yield 34 news-derived features per parameter combination.

3.4.6 Feature Integration & Scaling

The complete feature matrix integrates baseline features with news-derived features to produce the input representation for predictive modelling. Table XX summarises the complete feature set:

Feature Category	*Count*	*Description*
Price Features	2	Day-ahead price, spot price
Spread Features	1	Absolute spread
Load Features	1	System load
Lagged Features	4	24h and 168h lags for price and load
Temporal Features	5	Hour, day of week, day/week/month of year
Topic Features	14	Time-decayed counts per leaf topic
Embeddings Features	20	UMAP-reduced embedding dimensions

The news-derived features (34 dimensions) account for approximately 72% of the total feature space, reflecting the research focus on extracting predictive signals from textual information.

The dataset is partitioned chronologically to preserve the temporal ordering essential for time-series forecasting:

- **Training Set**: First 70% of observations
- **Validation Set**: Next 20% of observations
- **Test Set**: Final 10% of observations

This temporal separation ensures that all test-set observations occur strictly after all training observations, reflecting the realistic constraint that forecasting models cannot access future information. The validation set serves for hyperparameter tuning and time-

decay parameter selection, while the test set provides an unbiased estimate of out-of-sample performance SOURCE.

3.5 Model Training and Optimisation

This section details the hierarchical training methodology employed to transform engineered features into actionable trading signals. The framework implements a two-stage stacked ensemble architecture wherein an XGBoost classifier first learns complex interactions between news-derived features and market dynamics, followed by a LightGBM meta-learner that refines these predictions into long/short/hold trading decisions. This approach draws upon established stacking generalisation principles, wherein base-level predictions serve as meta-features for subsequent learners, improving overall predictive accuracy (Wolpert, 1992, pp. 241–259).

3.5.1 Training Framework Architecture

The training pipeline follows a systematic progression through five computational stages, designed to maintain strict temporal ordering whilst maximising parallel efficiency where dependencies permit. Building upon the pre-processed features described in Section 3.4, the workflow proceeds as follows.

The two-stage stacked architecture addresses a fundamental challenge in news-driven electricity forecasting: the need to extract predictive signal from high-dimensional, noisy textual features whilst avoiding overfitting on limited training samples. As discussed in Section 2.2.3, limitations of NLP in financial applications include the risk of spurious correlations when combining textual and numerical features. The hierarchical approach mitigates this by progressively reducing dimensionality across stages.

The first stage employs XGBoost as the base learner because it is effective in electricity price forecasting, where ensemble methods consistently outperform single-model alternatives (Lago et al., 2018, pp. 395–297). XGBoost's regularised gradient-boosting framework provides implicit feature selection that identifies which news-derived signals contain genuine predictive content. However, careful regularisation is required to prevent spurious correlations when combining BERT-based news features with price series (Guan et al., 2022, pp. 3573–3575).

The second stage employs LightGBM as the meta-learner, operating on a substantially reduced feature space comprising baseline price features and calibrated XGBoost probability outputs. This dimensionality reduction aims to mitigate the risk of overfitting whilst enabling the meta-learner to learn optimal weightings of the base model's predictions across different market conditions. The approach aligns with the stacking generalisation framework formalised by Wolpert (1992), in which base-level predictions serve as meta-features for subsequent learners, thereby improving overall predictive accuracy.

The full workflow proceeds through the following stages:

1. **Feature Parameter Selection**: Evaluate 30 time-decay parameter combinations using Ridge regression with cross-validation (CV), selecting the top five configurations for subsequent model training.
2. **XGBoost Hyperparameter Optimisation**: Conduct randomised search across 80 hyperparameter configurations for each selected parameter combination, employing expanding-window CV to preserve temporal dependencies.
3. **Probability Calibration**: Apply Platt scaling to transform XGBoost output scores into well-calibrated probability estimates suitable for meta-feature generation SOURCE.
4. **Meta-Learner Training**: Train LightGBM classifiers on both signal-enriched features (baseline + XGBoost predictions) and baseline-only features for comparative analysis.

The chronological data split established in Section 3.4.6 is maintained throughout, ensuring that all model development decisions are made without access to future observations.

3.5.2 Feature Parameter Selection via Ridge Regression

Before committing substantial computational resources to gradient boosting optimisation, the pipeline implements a lightweight screening phase using Ridge regression classifiers. This approach efficiently identifies promising time-decay parameter combinations from the 30-element grid (6 lookback windows × 5 decay rates) described in Section 3.4.4.

Ridge regression incorporates L2 regularisation to constrain coefficient magnitudes, mitigating multicollinearity in high-dimensional feature spaces (Hoerl & Kennard, n.d., pp. 55–58). The regularisation strength α is selected via grid search over 13 logarithmically spaced values spanning $[10^{-3}, 10^{3}]$. For each of the 30 parameter combinations, Ridge classifiers are fitted using 5-fold CV with accuracy as the scoring metric.

Following cross-validation, the five parameter combinations achieving the highest validation accuracy proceed to XGBoost training. This two-stage screening approach serves dual purposes: (i) reducing computational burden by eliminating poorly performing configurations early, and (ii) enabling comparative analysis of how different temporal aggregation strategies influence predictive performance. The efficiency of linear models for initial parameter screening has been demonstrated in similar ensemble pipelines (Caruana et al., 2004, pp. 3–5).

3.5.3 XGBoost Hyperparameter Optimisation

XGBoost implements regularised gradient boosting with built-in L1 and L2 penalties that control model complexity via tree-structure and leaf-weight constraints (T. Chen & Guestrin, 2016, pp. 785–789).

Rather than an exhaustive grid search, the pipeline employs randomised search sampling from continuous distributions. Following Bergstra et al. (2012, pp. 288–290), random search with 60–100 iterations over a 9-dimensional hyperparameter space yields results comparable to exhaustive grid search while reducing computational burden by orders of magnitude. Specifically, randomised search more efficiently explores high-dimensional parameter spaces by avoiding the combinatorial explosion inherent in grid-based approaches.

The search explores 80 random configurations per dataset, yielding 640 total model fits when combined with 8-fold CV. The hyperparameter distributions are configured to balance model complexity against regularisation strength. Table XX details the search distributions.

Parameter	Distribution	Range	Purpose
n_estimators	Uniform integer	[100, 400]	Number of boosting rounds

max_depth	Uniform integer	[2, 6]	Maximum tree depth
learning_rate	Log-uniform	[0.01, 0.15]	Step size shrinkage
subsample	Uniform	[0.5, 0.9]	Row sampling ratio
colsample_bytree	Uniform	[0.5, 0.9]	Column sampling ratio
gamma	Uniform	[0.5, 5.0]	Minimum loss reduction for split
min_child_weight	Uniform integer	[3, 12]	Minimum instance weight in child
reg_alpha	Log-uniform	[10^{-3}, 10^{1}]	L1 regularisation term
reg_lambda	Log-uniform	[10^{-2}, 10^{2}]	L2 regularisation term

Table XX: XGBoost Hyperparameter Search Distributions

The distributions are deliberately constrained relative to typical defaults: the maximum depth is limited to 6 (rather than the common ceiling of 10–12), the upper bound on the learning rate is reduced to 0.15 (from 0.3), and the lower bounds on the regularisation parameters are increased. These modifications address the validation-to-test degradation pattern characteristic of overfitting in time-series contexts with limited training samples.

To preserve the temporal ordering established in Section 3.4.6, the pipeline implements an expanding-window (walk-forward) validation scheme in which the training set grows monotonically whilst the test window advances in time (Tashman, 2000, pp. 440–445). The pipeline uses eight folds with a 72-hour step size and a 336-hour minimum training duration, balancing variance estimation against computational cost.

To address imbalances in the three-class target distribution, the pipeline uses squared inverse-frequency sample weighting, assigning observation weights proportional to the square of the inverse of class prevalence. This more aggressive weighting scheme amplifies the contribution of minority-class observations to the loss-function gradient, thereby encouraging the model to produce predictions for severely underrepresented classes rather than defaulting to majority-class predictions (Haibo He & Garcia, 2009, pp. 1270–1275).

To prevent overfitting during boosting iterations, early stopping monitors the validation set performance and terminates training when the evaluation metric no longer improves (Prechelt, 1998, pp. 761–763). The configuration employs:

- **Evaluation metric**: Multi-class log-loss (mlogloss)
- **Early stopping rounds**: 20 iterations without improvement
- **Evaluation set**: Held-out validation partition (20% of data)

Early stopping operates independently of CV, providing an additional regularisation mechanism that adapts the effective number of boosting rounds to each hyperparameter configuration. This dynamic stopping criterion reduces computational waste on overparameterised configurations whilst allowing well-regularised models to train longer.

For the three-class classification task, the XGBoost objective is configured as *multi:softprob*, which optimises multi-class log-loss whilst outputting probability estimates for each class. The macro-averaged F1 score serves as the primary optimisation target during hyperparameter search, balancing precision and recall across all three classes equally, regardless of class prevalence:

$$F_{1,\text{macro}} = \frac{1}{K}\sum_{c=1}^{K}\frac{2 \cdot \text{Precision}_c \cdot \text{Recall}_c}{\text{Precision}_c + \text{Recall}_c}$$

This choice reflects the trading application's requirement for reliable signals across all market conditions. Macro-averaging ensures that minority-class performance receives equal weight in model selection, preventing the optimiser from sacrificing Neutral-class accuracy to improve performance on the more prevalent Long and Short classes (Sokolova & Lapalme, 2009, pp. 430–431).

3.5.4 Probability Calibration via Platt Scaling

Whilst XGBoost with the *softprob* objective produces normalised probability vectors, these outputs may not be well-calibrated in the sense that a predicted probability of 0.7 corresponds to 70% empirical accuracy (Niculescu-Mizil & Caruana, n.d., pp. 2–4). Poorly calibrated probabilities can degrade downstream model performance when used as meta-features, as the LightGBM meta-learner may learn to compensate for systematic biases rather than exploiting genuine predictive signal.

Platt scaling fits a sigmoid function to transform classifier scores into calibrated probabilities, learning parameters via maximum likelihood on a held-out calibration set(Platt, 1999, pp. 63–66). This approach is well-suited for gradient-boosting ensembles, which often exhibit systematic over- or underconfidence (Niculescu-Mizil & Caruana, n.d., pp. 1–2).

The pipeline applies Platt scaling to the best-performing XGBoost model using the following configuration:

- **Method**: Sigmoid (Platt scaling, as opposed to isotonic regression)
- **Cross-validation mode**: *cv='prefit'*, to indicate the base model is already trained
- **Calibration set**: Held-out validation partition (20% of data)

Using the prefit mode applies calibration without retraining the base classifier, fitting the sigmoid parameters directly to validation-set predictions. This approach avoids the computational cost of refitting XGBoost whilst still correcting for probability miscalibration. The calibrated model replaces the uncalibrated version for all downstream operations, including meta-feature generation for LightGBM training.

3.5.5 Meta-Learner Training with LightGBM

LightGBM employs gradient-based sampling and feature bundling optimisations that achieve substantial speedups over traditional gradient boosted decision tree implementations whilst maintaining predictive accuracy (Ke et al., 2017).

The training framework implements stacked generalisation (Wolpert, 1992) by using XGBoost predictions as meta-features for LightGBM. In this two-level architecture, the base learner (XGBoost) generates predictions that serve as inputs to the meta-learner (LightGBM), thereby enabling the ensemble to learn complex combinations of base-model outputs. Specifically, four features are extracted from calibrated XGBoost output:

1. *xgb_pred:* Predicted class label (-1, 0, or 1)
2. *xgb_prob_class0:* Probability of Neutral class
3. *xgb_prob_class1:* Probability of Long class
4. *xgb_prob_class2:* Probability of Short class

These meta-features capture both the hard prediction and the confidence distribution, enabling the meta-learner to weight XGBoost contributions appropriately. Research on stacking ensembles demonstrates that probability outputs typically yield superior meta-features than class labels alone, as they encode uncertainty information that is lost in hard predictions (Džeroski & Ženko, 2004, pp. 266–268).

To quantify the contribution of news-derived signals, two LightGBM models are trained in parallel:

- **Signal Model**: Combines baseline price features with XGBoost meta-features to represent the full predictive pipeline, incorporating news-derived information.
- **Baseline Model**: Uses price features only, providing a reference for ablation analysis that isolates the marginal contribution of news signals.

Comparing these models provides direct insights for whether news-derived signals add predictive value beyond price history. If the signal model significantly outperforms the baseline, this validates the news feature engineering methodology; conversely, if performance is comparable, it suggests limited information content in the news aggregations.

Unlike the XGBoost randomised search, LightGBM employs exhaustive grid search over a more constrained parameter space, reflecting the meta-learner's role in refinement rather than primary feature learning. Exhaustive grid search is feasible for the meta-learner's reduced feature space (7 inputs vs. 47 for XGBoost), where the smaller search volume permits complete enumeration without the sampling variance inherent in random search. Table XX presents the parameter grid:

Table XX: LightGBM Grid Search Parameters

Parameter	***Values***	***Description***
num_leaves	[31, 50, 100]	Maximum leaves per tree
max_depth	[5, 10, 15]	Maximum tree depth
learning_rate	[0.01, 0.05, 0.1]	Boosting step size
n_estimators	[100]	Fixed boosting rounds
min_child_samples	[20]	Minimum samples per leaf
subsample	[0.8]	Row sampling ratio
colsample_bytree	[0.8]	Column sampling ratio

LightGBM training mirrors the XGBoost configuration with expanding-window CV:

- **Number of splits**: 5 folds, reduced from 8 to reflect the meta-learner's simpler optimisation landscape and the computational efficiency requirements during ablation analysis across multiple model variants

- **Step size**: 72 hours
- **Minimum training size**: 336 hours
- **Early stopping rounds**: 20 iterations without improvement

Class imbalance is addressed in a manner similar to XGBoost, using inverse-frequency sample weights computed from the training set's class prevalence (instead of square inverse-frequencies).

3.6 Model Evaluation Framework

TBA

TBA

3.7 Backtesting Framework

LOL

This section details the backtesting methodology employed to translate model predictions into simulated trading outcomes and evaluate economic performance. While Section 3.6 assessed predictive accuracy through classification metrics, this section evaluates the practical utility of the signal-generating pipeline by simulating trading execution on held-out test data and comparing performance against benchmark strategies.

3.7.1 Trading Strategy Specification

Purpose: Define how model predictions are translated into executable trading actions and establish the benchmark strategies for comparison.

- **Action mapping**: Model predictions (Long/Neutral/Short) are translated directly into trading positions:

  - **Long (+1)**: Buy at day-ahead auction, sell at intraday spot — profits when spot > day-ahead

- **Short (−1)**: Sell at day-ahead auction, buy at intraday spot — profits when spot < day-ahead

- **Neutral (0)**: No position taken — neither profit nor loss

- **Signal generation process** (two-stage pipeline):

1. XGBoost base model generates calibrated probability estimates for spread direction

2. LightGBM meta-learner refines these into final Long/Neutral/Short signals

- Note that the Neutral class serves a risk management function, allowing the model to abstain when prediction confidence is insufficient

- Baseline Model Strategy

- Uses the price-only LightGBM classifier (without news-derived features) as the baseline

- Same action mapping as the signal model

- Enables direct comparison of news feature contribution to trading performance

- Naïve Benchmark Strategy

- **Always-long strategy**: Systematically buys at day-ahead auction and sells at intraday spot for every observation

- Equivalent to setting action = +1 for all timestamps

- Represents a simple market-making approach that profits from average positive spreads but incurs losses during negative spread periods

- Provides a lower bound for model utility—if the signal model cannot outperform this naïve strategy, the predictive complexity adds no economic value

- Transaction Cost Modelling

- **Fixed costs**: Per-trade fee in EUR/MWh representing exchange fees, clearing costs, and broker commissions

- **Proportional costs**: Percentage of trade value representing bid-ask spreads and market impact

- Returns computation:

$$R_t = \begin{cases}

\text{spread}_t - c_{\text{fixed}} - c_{\%} \cdot |\text{spread}_t| & \text{if action} = +1 \\

-\text{spread}_t - c_{\text{fixed}} - c_{\%} \cdot |\text{spread}_t| & \text{if action} = -1 \\

0 & \text{if action} = 0

\end{cases}$$

- Note that transaction costs are only incurred when positions are taken (Long or Short), not for Neutral signals

Performance Metrics

Purpose: Define the economic performance metrics used to evaluate trading strategy outcomes.

- **Total Return**: Cumulative profit/loss over the test period in EUR/MWh

$$\text{Total Return} = \sum_{t=1}^{T} R_t$$

- **Average Return**: Mean return per period, indicating expected profit per trade

$$\bar{R} = \frac{1}{T} \sum_{t=1}^{T} R_t$$

- **Volatility**: Standard deviation of returns, measuring return dispersion

$$\sigma = \sqrt{\frac{1}{T-1} \sum_{t=1}^{T} (R_t - \bar{R})^2}$$

- **Sharpe Ratio** (annualised): Return per unit of total risk, the standard risk-adjusted performance measure (Sharpe, 1966)

$$\text{Sharpe} = \frac{\bar{R}}{\sigma} \cdot \sqrt{N_{\text{periods/year}}}$$

where $N_{\text{periods/year}} = 24 \times 365 = 8{,}760$ for hourly data

- **Sortino Ratio** (annualised): Return per unit of downside risk, penalising only negative volatility (Sortino & van der Meer, 1991)

$$\text{Sortino} = \frac{\bar{R}}{\sigma_{\text{downside}}} \cdot \sqrt{N_{\text{periods/year}}}$$

where σ_{downside} is computed only from negative returns

- Rationale: Upside volatility is desirable in trading; Sortino isolates harmful downside risk

- **Maximum Drawdown**: Largest peak-to-trough decline in cumulative returns

$$\text{MDD} = \min_t \left(\text{Cumulative}_t - \max_{s \leq t} \text{Cumulative}_s \right)$$

- Measures worst-case capital erosion during the strategy's lifetime

- Critical for assessing tail risk and capital requirements

- **Win Rate**: Percentage of profitable trades among all executed trades

$$\text{Win Rate} = \frac{\text{Number of } R_t > 0}{\text{Number of } R_t \neq 0} \times 100\%$$

- **Number of Trades**: Total count of non-zero positions (Long + Short signals)

- Lower trade counts with similar returns indicate more efficient signal generation

- Sharpe and Sortino ratios normalise for risk, enabling fair comparison

- Maximum drawdown indicates worst-case scenarios that may be unacceptable regardless of average returns

- Equity Curve Visualisation

- Plot cumulative returns over time for all strategies

- Enables visual identification of:

- Periods of sustained outperformance/underperformance

- Regime changes where model effectiveness varies

- Drawdown events and recovery patterns

- Sensitivity Analysis (Optional Extension)

- Evaluate performance under varying transaction cost assumptions

- Test robustness of strategy ranking to cost parameter changes

- Identify break-even transaction costs where model loses edge over naïve benchmark

4 Results

The subsequent chapter is devoted to the presentation of the empirical results of the analysis and the evaluation of the performance of the proposed news-augmented modelling framework. The results obtained are structured in accordance with the sequential stages of the research design.

Section 4.1 provides a comprehensive overview of the finalised dataset, accompanied by a detailed exposition of the descriptive statistics pertaining to the energy market variables, the news corpus, and the target variable employed for the classification process. In Section 4.2, the outcomes of the feature engineering process are examined, including the results of zero-shot topic classification, embedding diagnostics, and the selection of time-decay and dimensionality-reduction parameters. Section 4.3 evaluates the predictive performance of the proposed models by comparing news-augmented specifications with baseline benchmarks using statistical accuracy measures and formal significance tests. Finally, Section 4.4 assesses the economic relevance of the derived signals through backtesting, reporting risk-adjusted performance metrics and drawdown characteristics under realistic transaction cost assumptions.

4.1 Data Overview and Descriptive Statistics

The final dataset used for the empirical analysis consists of 141’716 news articles and 43’798 hourly observations of the German electricity market after all preprocessing, filtering, and alignment steps. These observations form the basis for the descriptive analysis and the subsequent evaluation of the forecasting and trading models.

Using a symmetric deadband of ±3.0 EUR/MWh to discretise the next-day spot–day-ahead price spread, the resulting class distribution is moderately imbalanced. Approximately 45.6% of observations are classified as Short, 43.3% as Long, and 11.1% as Neutral (Appendix 1). This distribution indicates that economically meaningful deviations between auction and intraday prices occur frequently, while small spreads within the deadband represent a minority of cases.

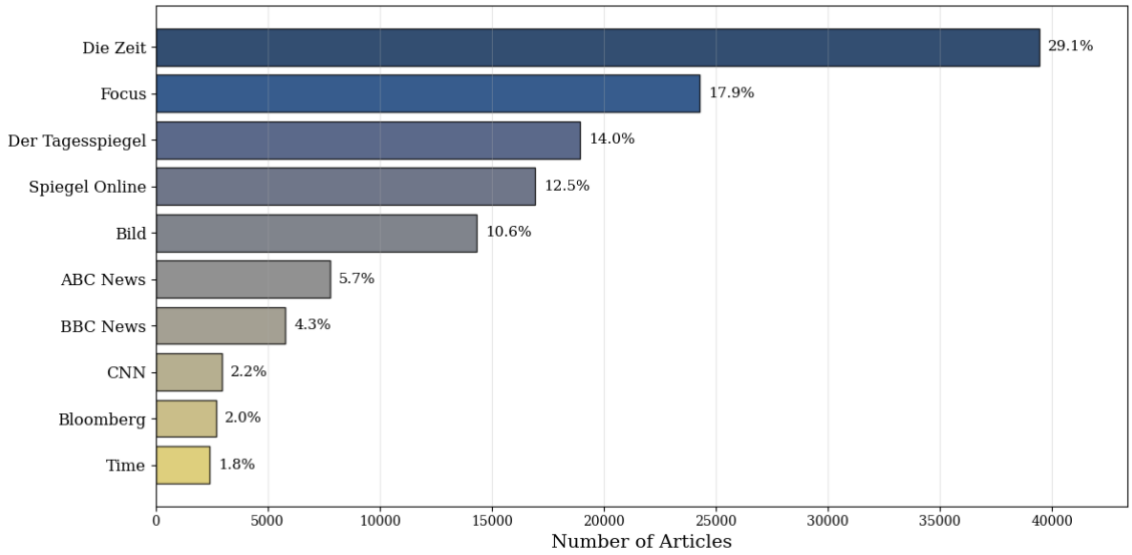

Figure 1: Top 10 News Sources

As illustrated in Figure XY, the corpus exhibits a high degree of concentration among primary news providers. The top five sources, comprising major German business-oriented outlets, account for 84% of all headlines. These are followed by five international sources (ranks six to eight) representing 16% of the data, while a long tail of smaller outlets provides supplementary coverage.

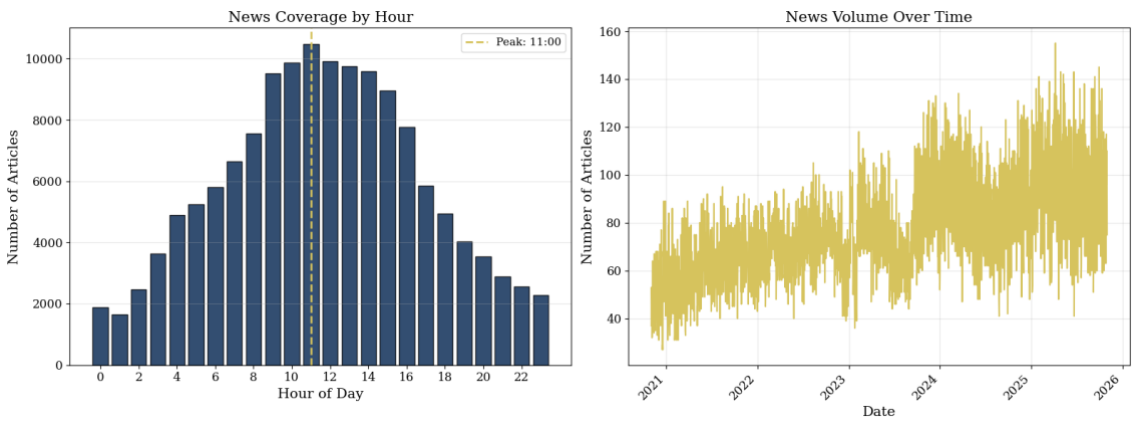

Figure 2: News Coverage by Hour of Day

The filtered news corpus shows pronounced temporal variation in publication intensity, as illustrated in Figure XY. News arrivals are concentrated during European business hours, with clear peaks in the late morning and early afternoon, and substantially lower activity during night-time hours. These temporal asymmetries imply that the informational environment faced by the model is uneven across hours, reinforcing the need for time-aware aggregation and decay mechanisms rather than uniform treatment of news events. In addition, the right panel of Figure XY shows a clear upward trend in

overall news volume over the sample period, reflecting an increasingly dense informational environment in energy markets, particularly in the later years of the dataset.

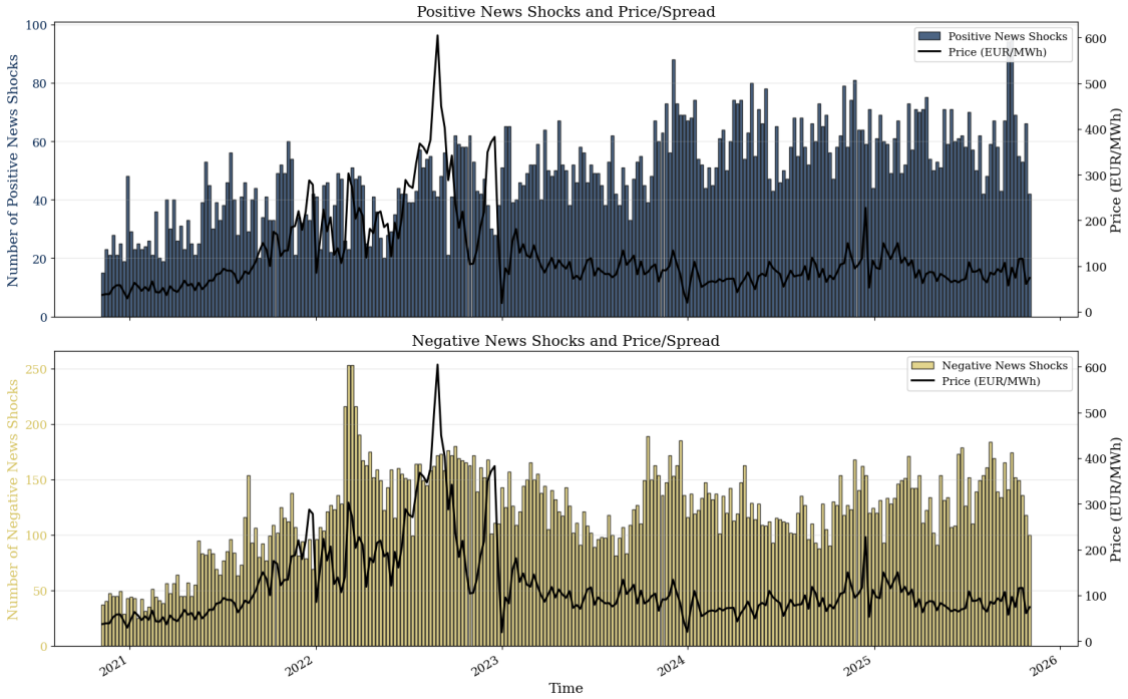

Figure 3: News Shocks and Price/Spread

Figure XX provides a dual-axis illustration of aggregated positive and negative news shocks, derived from the zero-shot topic classification, alongside the corresponding price or spread series. While periods of elevated news activity sometimes coincide with notable price movements, the high level of aggregation and heterogeneity of events limit the insights that can be drawn from visual inspection alone. Accordingly, the figure is presented for descriptive purposes and serves to motivate the model-based analysis that follows.

Taken together, the descriptive statistics and visual diagnostics highlight four key characteristics of the data: concentration in news sources, heterogeneity in topical coverage, strong temporal clustering of news arrivals, and time-varying alignment between news shocks and price movements. These characteristics provide essential context for the feature engineering outcomes and model performance results discussed in Sections 4.2 to 4.4.

4.2 Feature Engineering Results

This section reports the empirical outcomes of the feature engineering pipeline described in Chapter 3. The results focus on the structure and composition of the news-derived

features, the quality of the resulting textual representations, and the selection of key aggregation and dimensionality parameters used for downstream modelling.

The zero-shot classification framework yields a structured decomposition of the news corpus into 12 energy-relevant topic categories, while approximately 23.0% of articles are classified as non-relevant ("kein Bezug zu Energie, Wetter oder Finanzmärkten") and excluded from further aggregation. This filtering outcome confirms that a substantial share of keyword-retrieved news content does not convey information with plausible relevance for electricity price dynamics. The retained articles form the basis for all subsequent topic-based and embedding-based features.

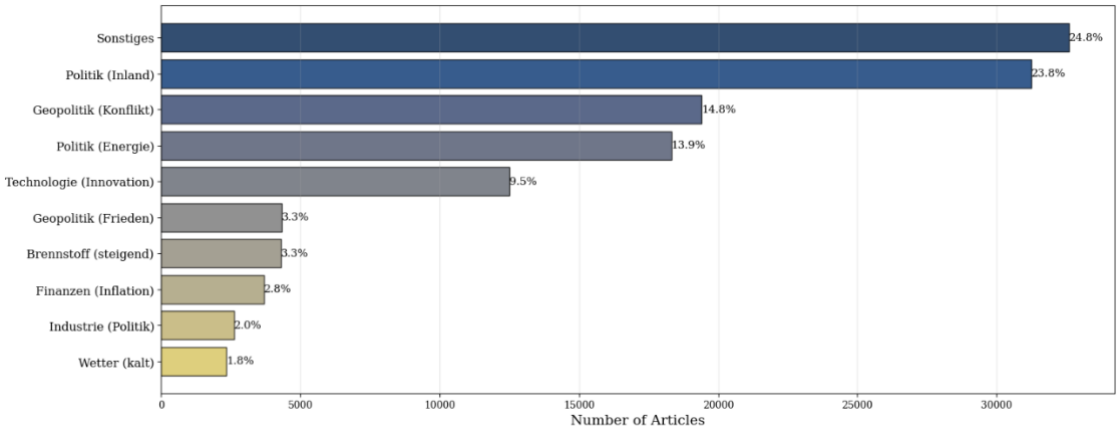

Figure 4: Top 10 News Topic Classifications

Figure XY summarises the topical structure of the filtered news corpus by displaying the ten most frequent topic classifications obtained from the zero-shot framework. Topics related to domestic political developments, general political and societal news, and geopolitical conflict dominate the corpus, followed by energy-related policy topics and innovation- or technology-oriented coverage. This distribution is consistent with the broader set of political, regulatory, and macroeconomic drivers of electricity price formation discussed in Chapter 2. At the same time, the distribution highlights the heterogeneity of the incoming news flow, even after relevance filtering. A more detailed breakdown of topic frequencies is provided in Appendix 2.

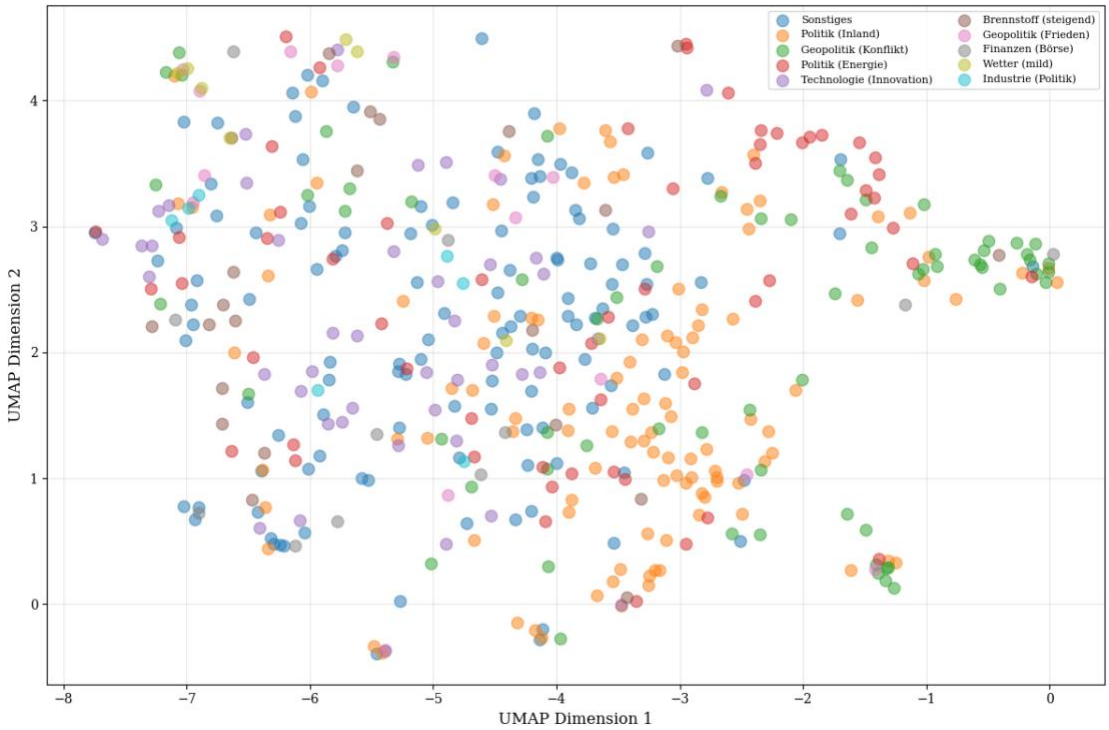

Figure 5: News Embedding Visualization (UMAP)

Beyond discrete topic labels, semantic structure in the news corpus is captured through dense sentence embeddings. Figure XX presents UMAP-based visual diagnostics of these embeddings, showing that articles assigned to similar topic categories cluster in embedding space while remaining distinguishable from unrelated categories. These projections suggest that the embedding representations preserve meaningful semantic structure and provide a complementary, continuous representation of news content beyond the categorical topic assignments.

UMAP DOES NOT FIX RANDOM SEEDS WHEN PROCESSING DATA IN PARALLEL

Table 2: Ridge Results

Dataset Name	Lookback Window in h	Decay-λ	Regularization Strength (α)	Validation Accuracy	Validation Macro F1
dataset_lw168_dl0.01	168	0.01	1000.0	0.516	0.362
dataset_lw504_dl0.01	504	0.01	1000.0	0.513	0.354
dataset_lw168_dl0.1	168	0.10	1000.0	0.512	0.354
dataset_lw504_dl0.1	504	0.10	100.0	0.510	0.357
dataset_lw48_dl0.5	48	0.50	1000.0	0.509	0.358

The temporal aggregation of news-derived features depends critically on the choice of time-decay parameters. A grid search over alternative lookback windows and decay rates

was evaluated using a Ridge classifier on the validation set. Table XX summarises the top-performing parameter combinations ranked by validation accuracy and macro-averaged F1 score. Across the tested configurations, models with relatively short to medium lookback windows and moderate decay rates achieve the strongest validation performance. In particular, the selected reference configuration with a 168-hour lookback window and decay parameter $\lambda = 0.01$ attains a validation accuracy of 0.516 and a macro-averaged F1 score of 0.362. This result indicates that news information retains predictive relevance over several days when incorporated with a relatively slow temporal decay. The identified parameterisation is therefore used as a benchmark for the subsequent evaluation of non-linear models.

To mitigate the risk of overfitting and reduce computational complexity, the original 384-dimensional sentence embeddings are compressed using UMAP into 20 lower-dimensional components. This reduction preserves the dominant semantic structure of the news corpus while substantially decreasing feature dimensionality. The resulting embedding components, together with the aggregated topic-count features, form the final set of news-derived inputs used in the predictive models evaluated in Sections 4.3 and 4.4.

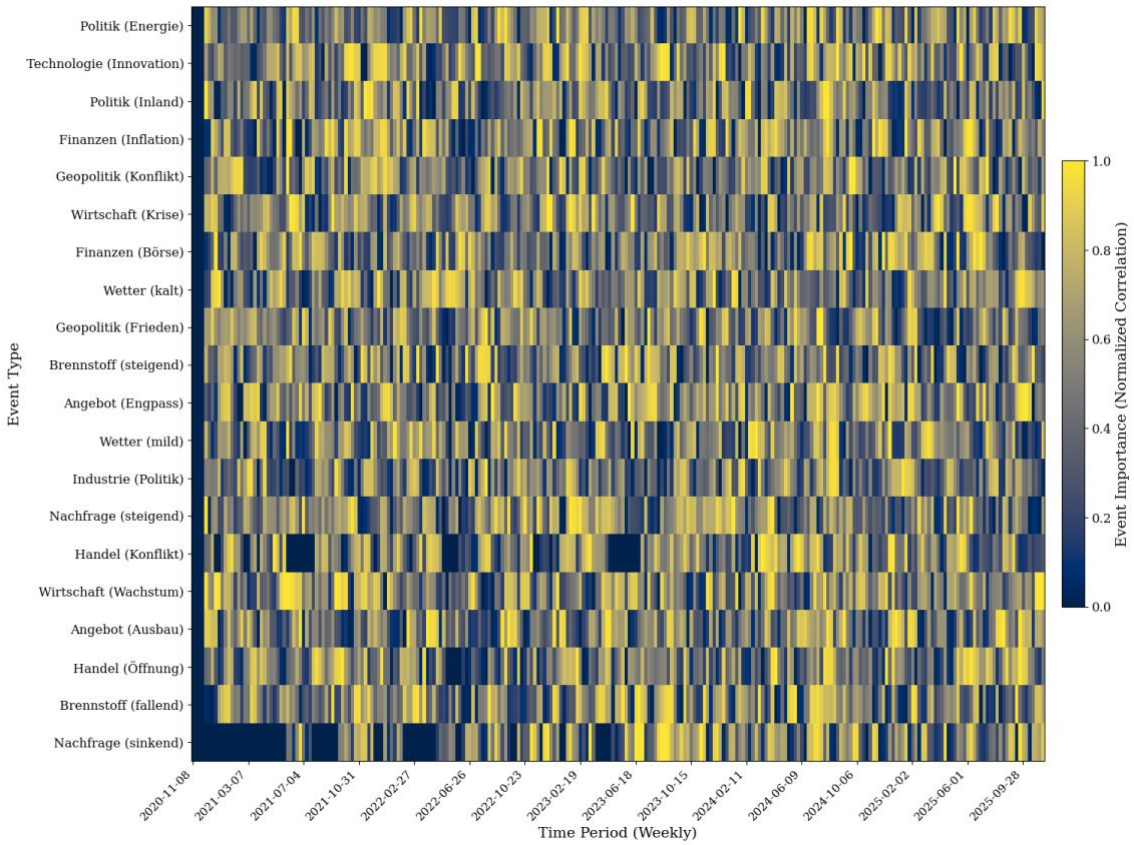

Figure 6: Variation in Event Importance Over Time

Figure XX illustrates temporal variation in the relevance of different news categories through an event-importance heatmap showing the correlation between topic-specific news intensity and subsequent spread movements across time periods. The heatmap indicates that the relationship between news topics and price dynamics is not constant over time. No single news category exhibits a persistent correlation with price movements across the sample. Instead, the relative importance of different event types varies across periods.

Taken together, these results demonstrate that the feature engineering pipeline yields a structured and semantically coherent representation of news information, while also revealing substantial time variation in topical relevance. These properties motivate the use of flexible, non-linear models and selective signal construction in the performance evaluation that follows.

4.3 Predictive Performance and Statistical Evaluation

This section evaluates the out-of-sample predictive performance of the proposed models and assesses whether the inclusion of news-derived information leads to statistically meaningful improvements over price-based benchmarks. All results are reported on the held-out test set and focus on multi-class classification performance across Long, Neutral, and Short signals.

The strongest standalone news-aware classifier is the XGBoost three-class model, trained on the combined set of market-based and news-derived features. Based on the training and validation results reported in Table X, the configuration with a 168-hour lookback window and a decay parameter of $\lambda = 0.01$ is selected for out-of-sample evaluation. On the held-out test set, this model achieves a test accuracy of 52.7%, a macro-averaged F1 score of 0.375, and a macro-averaged AUC of 0.61, with detailed results reported in Appendix 3. Learning curve diagnostics (Appendix 4) further indicate a gradual improvement in validation performance as the training sample expands, suggesting that additional data could lead to higher out-of-sample accuracy. These results indicate that the model captures economically relevant information beyond random classification in a setting characterised by class imbalance and noisy short-term price dynamics.

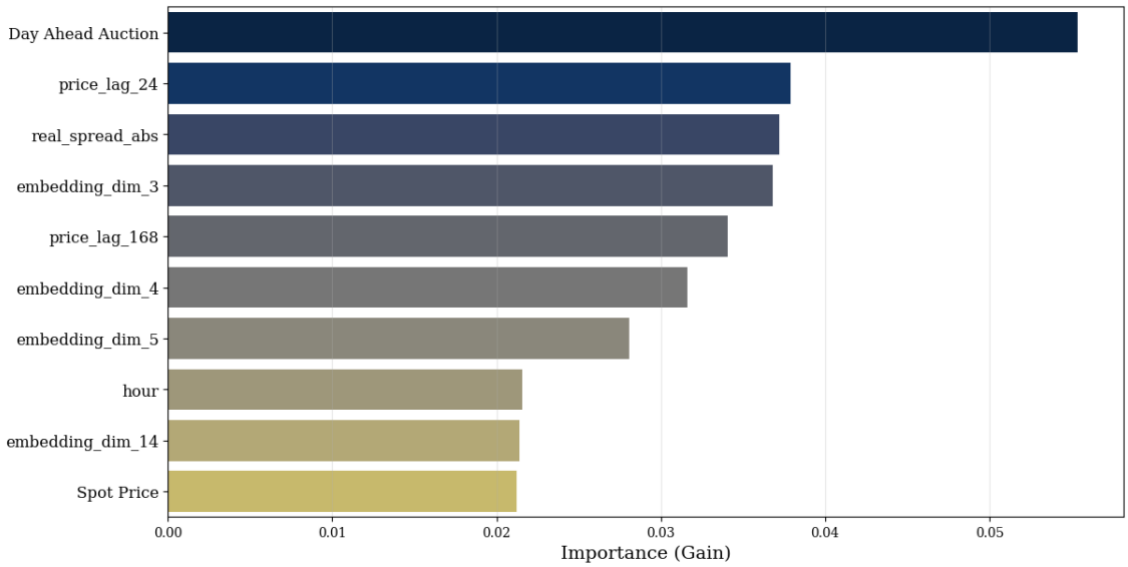

Figure 7: XGBoost Top 10 Features (Gain)

Figure XX shows the top ten features by gain-based importance in the XGBoost classifier. Market-based variables, such as the day-ahead auction price and price lags, dominate the

ranking, reflecting strong persistence in electricity prices. At the same time, several news-derived embedding dimensions appear among the most influential features, indicating that textual information provides incremental predictive value beyond historical prices alone.

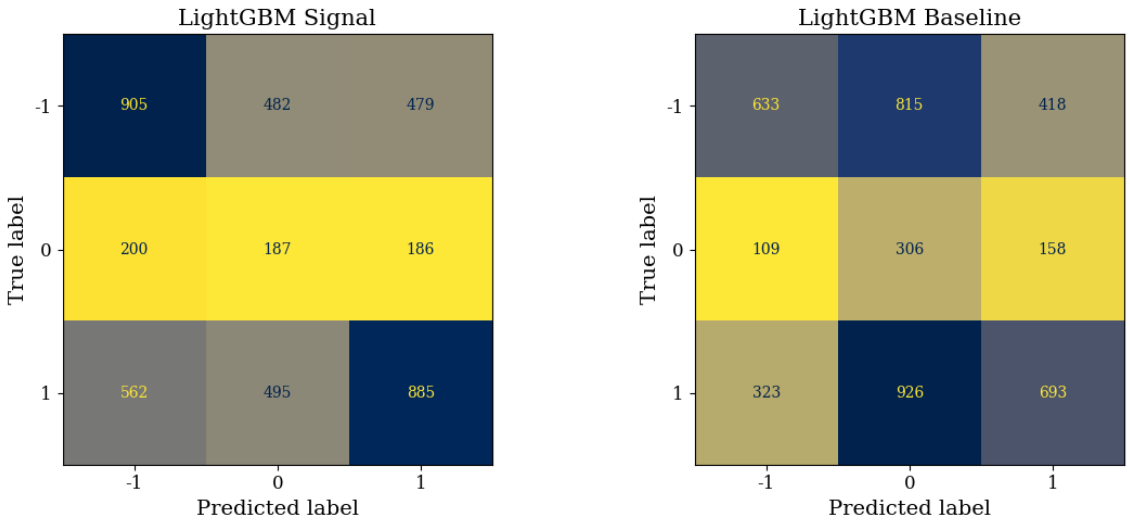

Figure 8: Confusion Matrices for Signal and Baseline Model

Table 3: LightGBM Classification Report (Test Set)

Class	Model	Precision	Recall	F1-Score	Support
Short (-1)	Signal	0.54	0.48	0.51	1866
	Baseline	0.59	0.34	0.43	1866
Neutral (0)	Signal	0.16	0.33	0.22	573
	Baseline	0.15	0.53	0.23	573
Long (1)	Signal	0.57	0.46	0.51	1942
	Baseline	0.55	0.36	0.43	1942
Macro Avg	Signal	0.42	0.42	0.411	4381
	Baseline	0.43	0.41	0.366	4381

	Signal	Baseline
Accuracy	0.451	0.373
Macro AUC	0.600	0.594

To translate the XGBoost forecasts into actionable trading signals, the predicted class probabilities are subsequently used as inputs to a second-stage LightGBM signal model. This model combines the XGBoost prediction outputs with a small set of baseline price features. As reported in Table X, the signal model achieves a test accuracy of 45.1% and

a macro-averaged F1 score of 0.411, outperforming the corresponding LightGBM baseline model, which relies solely on price-based inputs and attains an accuracy of 37.3% and a macro-F1 score of 0.366. While overall accuracy is lower than that of the standalone XGBoost classifier, the signal model exhibits a more balanced performance across classes, reflected in its higher macro-F1 score. Figure X presents the corresponding confusion matrices for the signal and baseline models, illustrating improved discrimination between Long and Short outcomes under the signal specification, while misclassification remains most pronounced for the Neutral class. The signal model also attains a slightly higher macro-averaged AUC than the baseline model, indicating improved class discrimination across probability thresholds; the corresponding ROC curves are shown in Appendix 5.

The performance difference between the signal model and the baseline model is formally assessed using McNemar's test for paired classification outcomes (Appendix 6). The resulting test statistic of $\chi^2 = 113.89$ with a corresponding p-value below 0.001 rejects the null hypothesis of equal predictive accuracy. Complementary bootstrap analysis of test-set accuracy further supports this finding, as the 95% confidence intervals of the two models do not overlap, providing robust evidence that the signal model delivers statistically significant improvements over the baseline.

An analysis of feature importance reveals a clear distinction between the two LightGBM specifications. In the signal model, the most influential features are the XGBoost prediction probabilities, indicating that the upstream classifier effectively condenses information from the news-derived inputs into a small number of high-value signals. In contrast, the baseline model relies primarily on price-based variables, such as the day-ahead auction price and absolute spread measures, reflecting its dependence on short-term price persistence rather than external information. The feature importance rankings for both models are reported in Appendix 7 and Appendix 8.

In summary, the empirical results show that models incorporating news-derived information achieve higher predictive performance than price-only benchmarks across multiple evaluation metrics. The signal model consistently attains higher macro-averaged F1 scores and AUC values and outperforms the baseline model in terms of predictive accuracy, with the difference confirmed as statistically significant by McNemar's test and bootstrap confidence intervals. At the same time, absolute performance levels remain

moderate, reflecting the noisy and volatile nature of short-term electricity price movements. These results complete the statistical evaluation of the predictive models.

4.4 Economic Performance and Backtesting Results

This section evaluates the economic performance of the proposed trading strategies based on the predictive signals examined in Section 4.3. A stylised portfolio backtest is conducted over the test period to assess whether the statistical improvements of the news-based signal model translate into superior trading outcomes relative to a price-based baseline and a naive benchmark strategy.

The portfolio simulation assumes an initial capital of €100'000 and applies a fixed position size of 10% of available capital per trade. Positions are taken on an hourly basis following the model-generated Long and Short signals, while Neutral signals result in no exposure. Transaction costs are modelled conservatively and consist of a fixed €0.50 per MWh and a proportional cost of 0.05% per trade. All strategies are evaluated under identical execution and cost assumptions to ensure comparability.

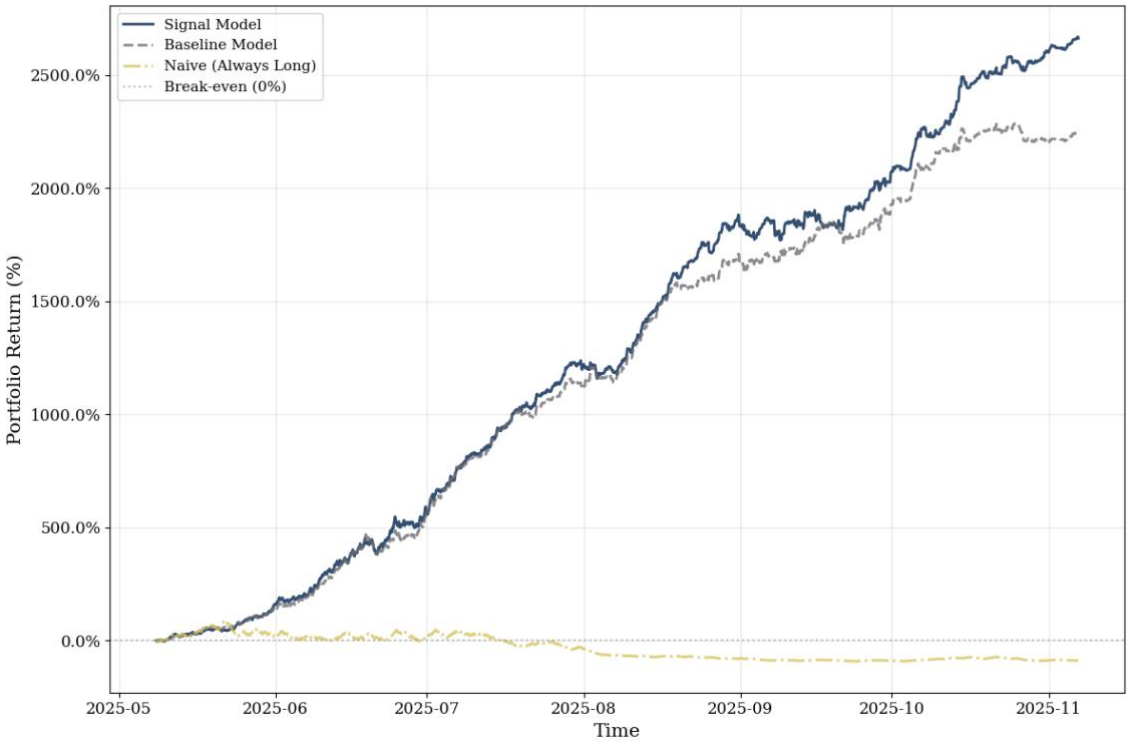

Figure 9: Portfolio Equity Curve - Test Period

Figure X displays the cumulative portfolio returns and corresponding drawdowns for the signal model, the baseline model, and a naive benchmark strategy that remains

continuously long. Both model-based strategies generate substantial cumulative gains over the test period, while the naive strategy experiences a severe and persistent drawdown, ultimately resulting in large capital losses. The signal model achieves the highest terminal portfolio value, reaching approximately €2'76 million, compared with €2'35 million for the baseline model. In contrast, the naive strategy ends the sample with a final portfolio value of roughly €13'000, reflecting the pronounced downside risk of an unconditional long exposure in the German power market.

Strategy	Final Value (EUR)	Total Return (%)	Annualized Return (%)	Sharpe Ratio	Sortino Ratio
Signal Model	168	0.01	1000.0	0.516	0.362
Baseline Model	504	0.01	1000.0	0.513	0.354
Naïve (Always Long)	168	0.10	1000.0	0.512	0.354

Table X summarises the key performance metrics for all three strategies. The signal model records a total return of 2.660%, outperforming the baseline model's 2.249% return. Risk-adjusted performance is substantially higher for the signal model, with an annualised Sharpe ratio of 1.552 and a Sortino ratio of 2.197, compared to 1.059 and 1.561, respectively, for the baseline strategy. Both models achieve win rates above 60%, while the naive strategy remains close to a coin-flip outcome.

The signal model executes a larger number of trades than the baseline strategy, reflecting its more active use of information, and incurs higher absolute transaction costs. Nevertheless, these costs are more than offset by the higher gross returns achieved. The naive strategy, despite executing the largest number of trades, generates strongly negative performance and exhibits the weakest risk-adjusted metrics.

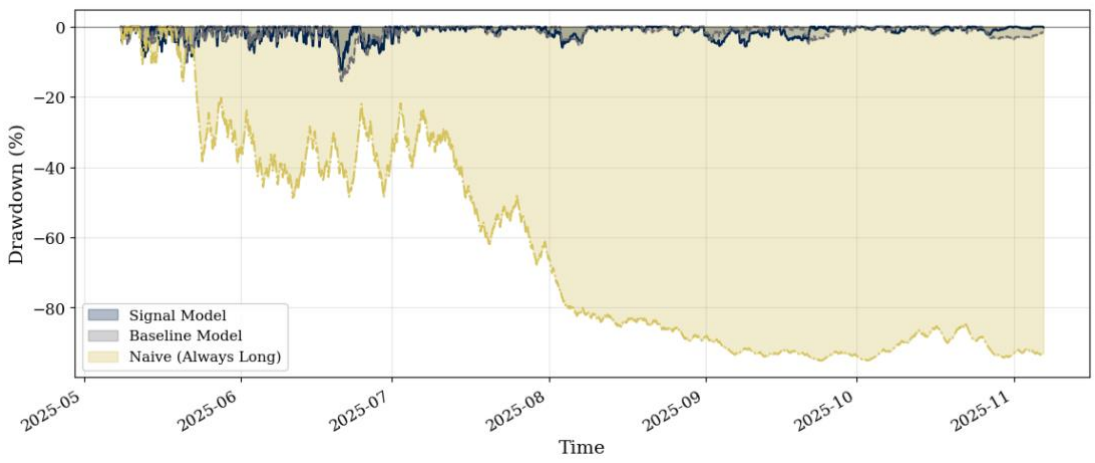

Figure 10: Drawdown (% of Peak)

In terms of trade-level outcomes, both model-based strategies achieve win rates above 60%, as reported in Table X, indicating that profitable trades occur more frequently than losing trades under both specifications. The baseline model records a slightly higher win rate than the signal model; however, this difference is accompanied by lower cumulative returns and weaker risk-adjusted performance.

Drawdown behaviour further differentiates the strategies. As summarised in Table X and illustrated in Figure X, the signal and baseline models both experience relatively shallow maximum drawdowns of approximately −4%, with drawdowns remaining contained throughout the test period. In contrast, the naive always-long strategy exhibits severe and persistent drawdowns, reaching a maximum drawdown of over −95%, with limited recovery. The drawdown trajectories in Figure X show that losses for the model-based strategies are short-lived and recover quickly, whereas the naive strategy remains exposed to prolonged downside risk.

Answer Research Question

Short Summary

5 Discussion

6 Conclusion

7 Bibliography

8 Appendix

Appendix 1: Target Distribution

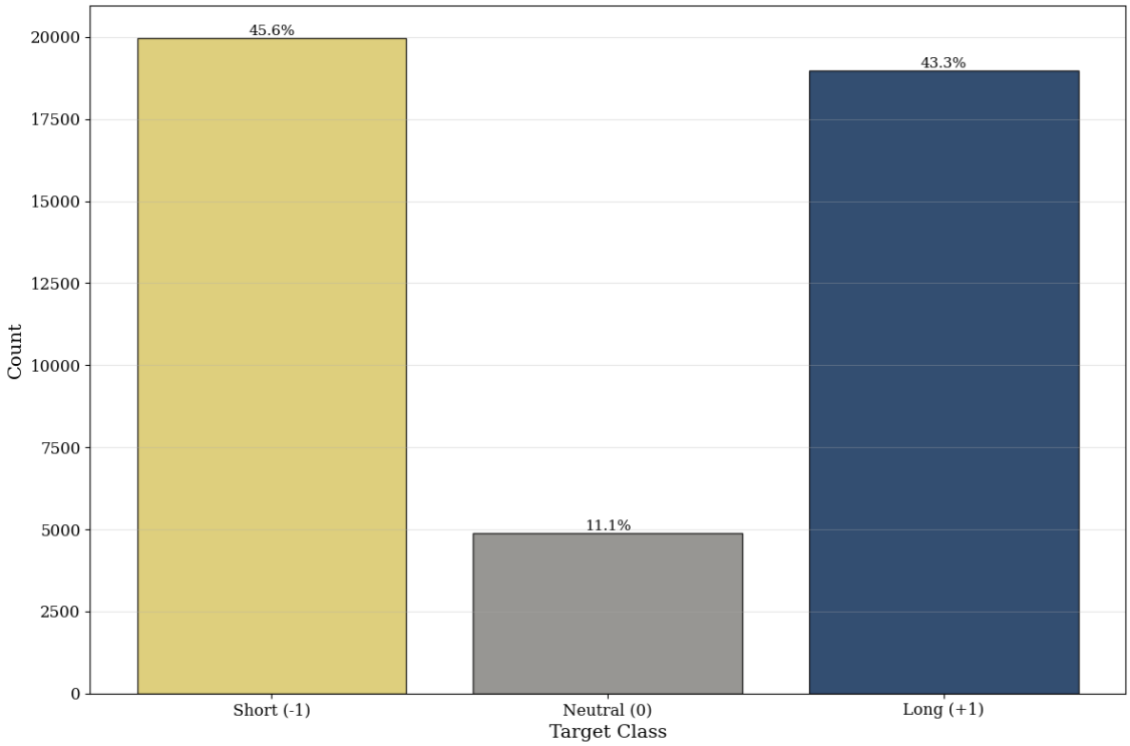

Appendix 2: News Distribution per Topic Classification

Table

Appendix 3: XGBoost Classification Report

Class	Precision	Recall	F1-Score	Support
Short (-1)	0.52	0.63	0.57	1866
Neutral (0)	0.16	0.11	0.13	573
Long (1)	0.56	0.50	0.53	1942

Accuracy	Macro-F1	Macro-Averaged AUC
0.5049	0.4103	0.610

Appendix 4: Learning curve diagnostics

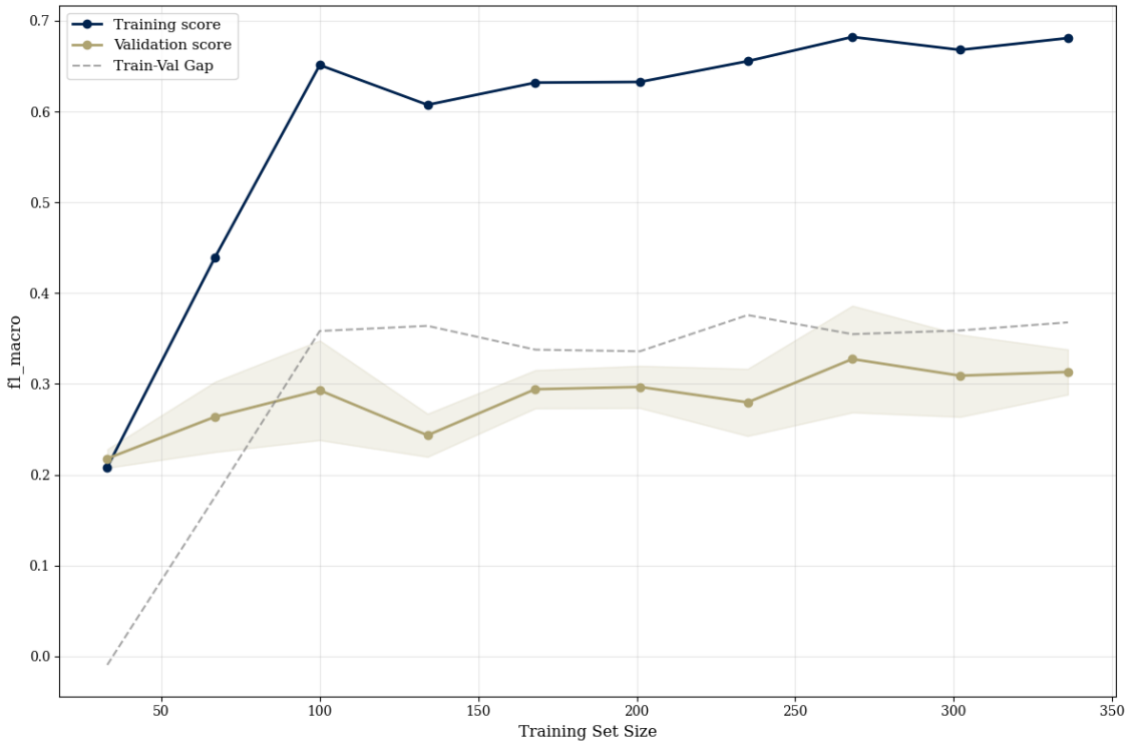

Appendix 5: ROC Curves for the LightGBM signal and baseline models (test set)

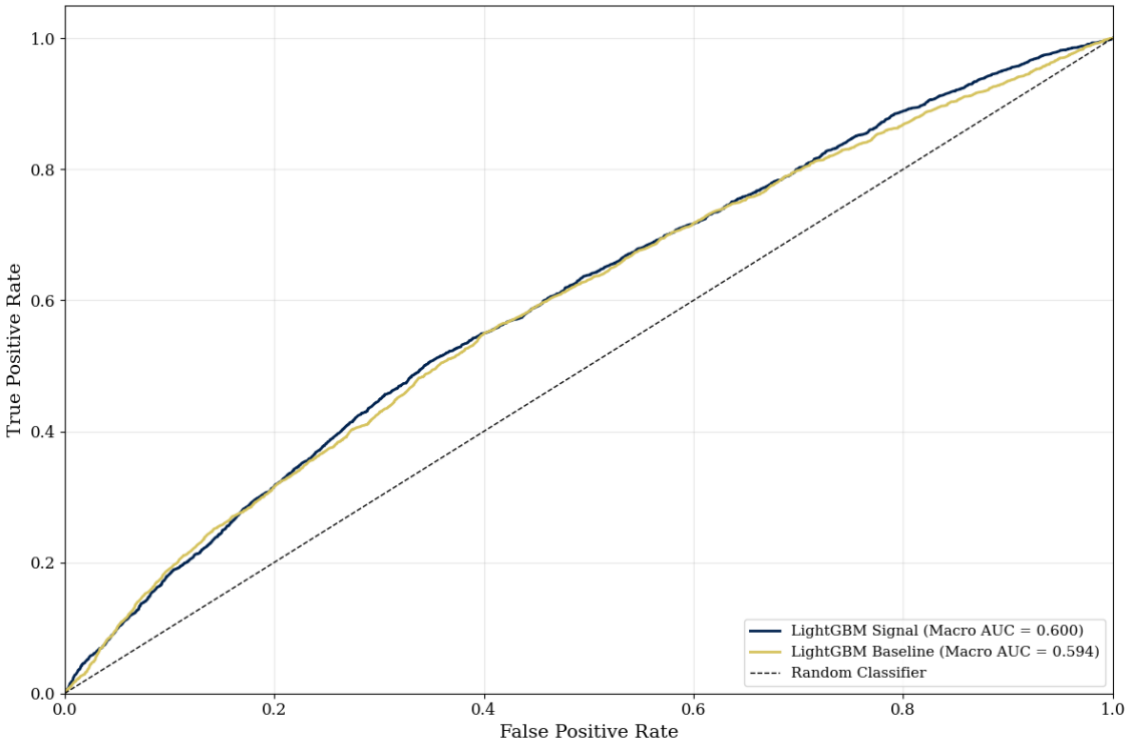

Appendix 6: Statistical Model Comparison with McNemar & Confidence Intervals

Panel A: Overall Predictive Accuracy

Model	Accuracy
Signal Model	0.4513
Baseline Model	0.3725

Panel B: McNemar's Test for Paired Comparison

Outcome Category	Count
Both correct	1285
Signal only	692
Baseline only	347
Both wrong	2057
Test Statistic	**Value**
McNemar's χ^2	113.8941
p-value	0.0000

Panel C: Bootstrap 95% Confidence Intervals for Accuracy

Signal model	0.4512 [0.4360, 0.4663]
Baseline model	0.3726 [0.3579, 0.3862]
Difference	+0.0786

Appendix 7: Features of LightGBM Signal Model

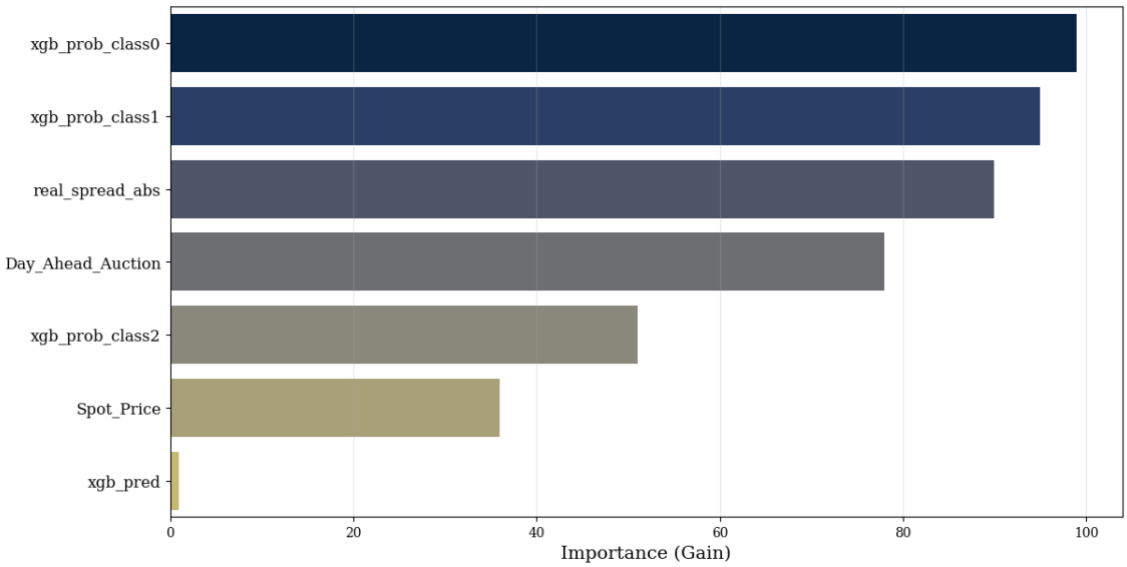

Appendix 8: Features of LightGBM Baseline Model

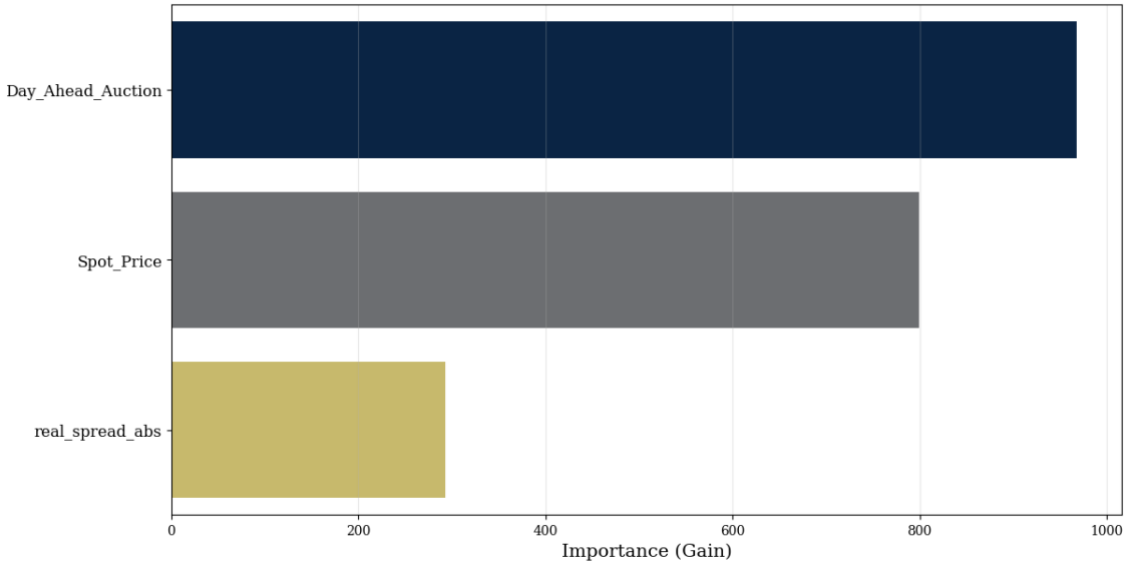

Appendix 1: Declaration of Independence

Appendix 2: AI Declaration

Appendix 3: Python Codes

Appendix 4: CSV File with all News Data

Appendix 5: CSV File with all Price and Market relevant Data